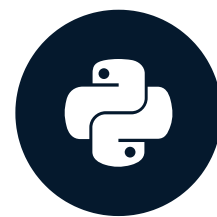


Rolling window functions with pandas

MANIPULATING TIME SERIES DATA IN PYTHON



Stefan Jansen

Founder & Lead Data Scientist at
Applied Artificial Intelligence

Window functions in pandas

- Windows identify sub periods of your time series
- Calculate metrics for sub periods inside the window
- Create a new time series of metrics
- Two types of windows:
 - Rolling: same size, sliding (this video)
 - Expanding: contain all prior values (next video)

Calculating a rolling average

```
data = pd.read_csv('google.csv', parse_dates=['date'], index_col='date')
```

```
DatetimeIndex: 1761 entries, 2010-01-04 to 2016-12-30  
Data columns (total 1 columns):  
price      1761 non-null float64  
dtypes: float64(1)
```



Calculating a rolling average

```
# Integer-based window size  
data.rolling(window=30).mean() # fixed # observations
```

```
DatetimeIndex: 1761 entries, 2010-01-04 to 2017-05-24  
Data columns (total 1 columns):  
price      1732 non-null float64  
dtypes: float64(1)
```

- `window=30` : # business days
- `min_periods` : choose value < 30 to get results for first days

Calculating a rolling average


```
# Offset-based window size  
data.rolling(window='30D').mean() # fixed period length
```

```
DatetimeIndex: 1761 entries, 2010-01-04 to 2017-05-24  
Data columns (total 1 columns):  
price      1761 non-null float64  
dtypes: float64(1)
```

- `30D` : # calendar days

90 day rolling mean

```
r90 = data.rolling(window='90D').mean()  
google.join(r90.add_suffix('_mean_90')).plot()
```



.join:
**concatenate Series or
DataFrame along
axis=1**


90 & 360 day rolling means

```
data['mean90'] = r90  
r360 = data['price'].rolling(window='360D').mean()  
data['mean360'] = r360; data.plot()
```




Multiple rolling metrics (1)

```
r = data.price.rolling('90D').agg(['mean', 'std'])  
r.plot(subplots = True)
```



Multiple rolling metrics (2)

```
rolling = data.google.rolling('360D')
q10 = rolling.quantile(0.1).to_frame('q10')
median = rolling.median().to_frame('median')
q90 = rolling.quantile(0.9).to_frame('q90')
pd.concat([q10, median, q90], axis=1).plot()
```



Let's practice!

MANIPULATING TIME SERIES DATA IN PYTHON

Expanding window functions with pandas

MANIPULATING TIME SERIES DATA IN PYTHON



Stefan Jansen

Founder & Lead Data Scientist at
Applied Artificial Intelligence

Expanding windows in pandas

- From rolling to expanding windows
- Calculate metrics for periods up to current date
- New time series reflects all historical values
- Useful for running rate of return, running min/max
- Two options with pandas:
 - `.expanding()` - just like `.rolling()`
 - `.cumsum()` , `.cumprod()` , `cummin()` / `max()`

The basic idea

```
df = pd.DataFrame({'data': range(5)})  
df['expanding sum'] = df.data.expanding().sum()  
df['cumulative sum'] = df.data.cumsum()  
df
```

	data	expanding sum	cumulative sum
0	0	0.0	0
1	1	1.0	1
2	2	3.0	3
3	3	6.0	6
4	4	10.0	10

Get data for the S&P 500

```
data = pd.read_csv('sp500.csv', parse_dates=['date'], index_col='date')
```

```
DatetimeIndex: 2519 entries, 2007-05-24 to 2017-05-24
```

```
Data columns (total 1 columns):
```

```
SP500      2519 non-null float64
```



How to calculate a running return

- Single period return r_t : current price over last price minus 1:

$$r_t = \frac{P_t}{P_{t-1}} - 1$$


- Multi-period return: product of $(1 + r_t)$ for all periods, minus 1:

$$R_T = (1 + r_1)(1 + r_2) \dots (1 + r_T) - 1$$

- For the period return: `.pct_change()`
- For basic math `.add()` , `.sub()` , `.mul()` , `.div()`
- For cumulative product: `.cumprod()`

Running rate of return in practice

```
pr = data.SP500.pct_change() # period return
pr_plus_one = pr.add(1)
cumulative_return = pr_plus_one.cumprod().sub(1)
cumulative_return.mul(100).plot()
```



Getting the running min & max

```
data['running_min'] = data.SP500.expanding().min()  
data['running_max'] = data.SP500.expanding().max()  
data.plot()
```




Rolling annual rate of return

```
def multi_period_return(period_returns):  
    return np.prod(period_returns + 1) - 1  
  
pr = data.SP500.pct_change() # period return  
r = pr.rolling('360D').apply(multi_period_return)  
data['Rolling 1yr Return'] = r.mul(100)  
data.plot(subplots=True)
```

Rolling annual rate of return

```
data['Rolling 1yr Return'] = r.mul(100)  
data.plot(subplots=True)
```



Let's practice!

MANIPULATING TIME SERIES DATA IN PYTHON

Case study: S&P500 price simulation

MANIPULATING TIME SERIES DATA IN PYTHON



Stefan Jansen

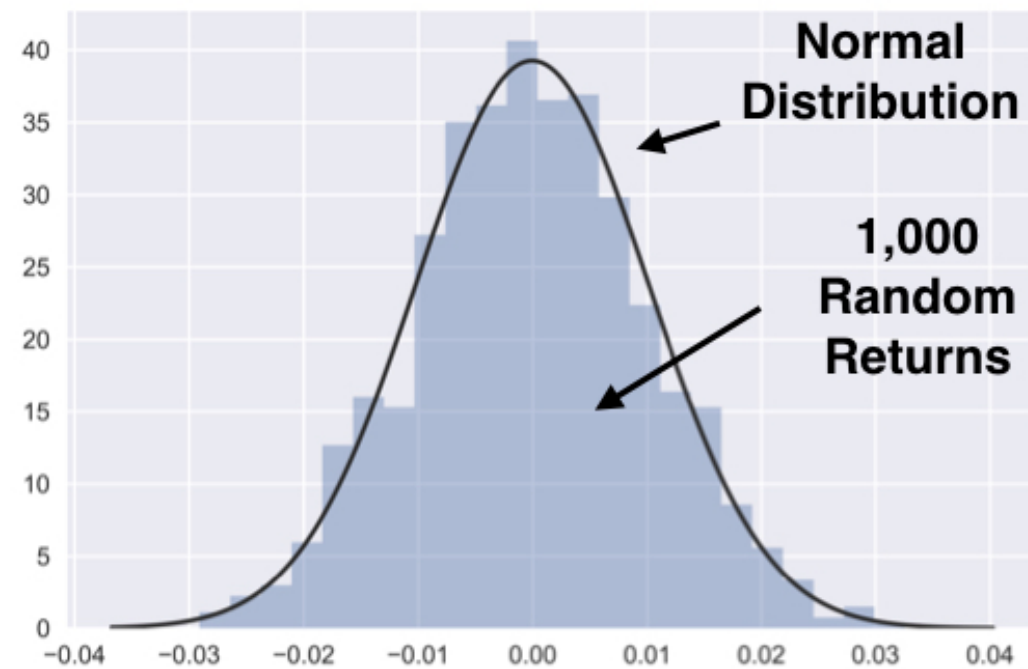
Founder & Lead Data Scientist at
Applied Artificial Intelligence

Random walks & simulations

- Daily stock returns are hard to predict
- Models often assume they are random in nature
- Numpy allows you to generate random numbers
- From random returns to prices: use `.cumprod()`
- Two examples:
 - Generate random returns
 - Randomly selected actual SP500 returns

Generate random numbers

```
from numpy.random import normal, seed
from scipy.stats import norm
seed(42)
random_returns = normal(loc=0, scale=0.01, size=1000)
sns.distplot(random_returns, fit=norm, kde=False)
```




Create a random price path

```
return_series = pd.Series(random_returns)
random_prices = return_series.add(1).cumprod().sub(1)
random_prices.mul(100).plot()
```



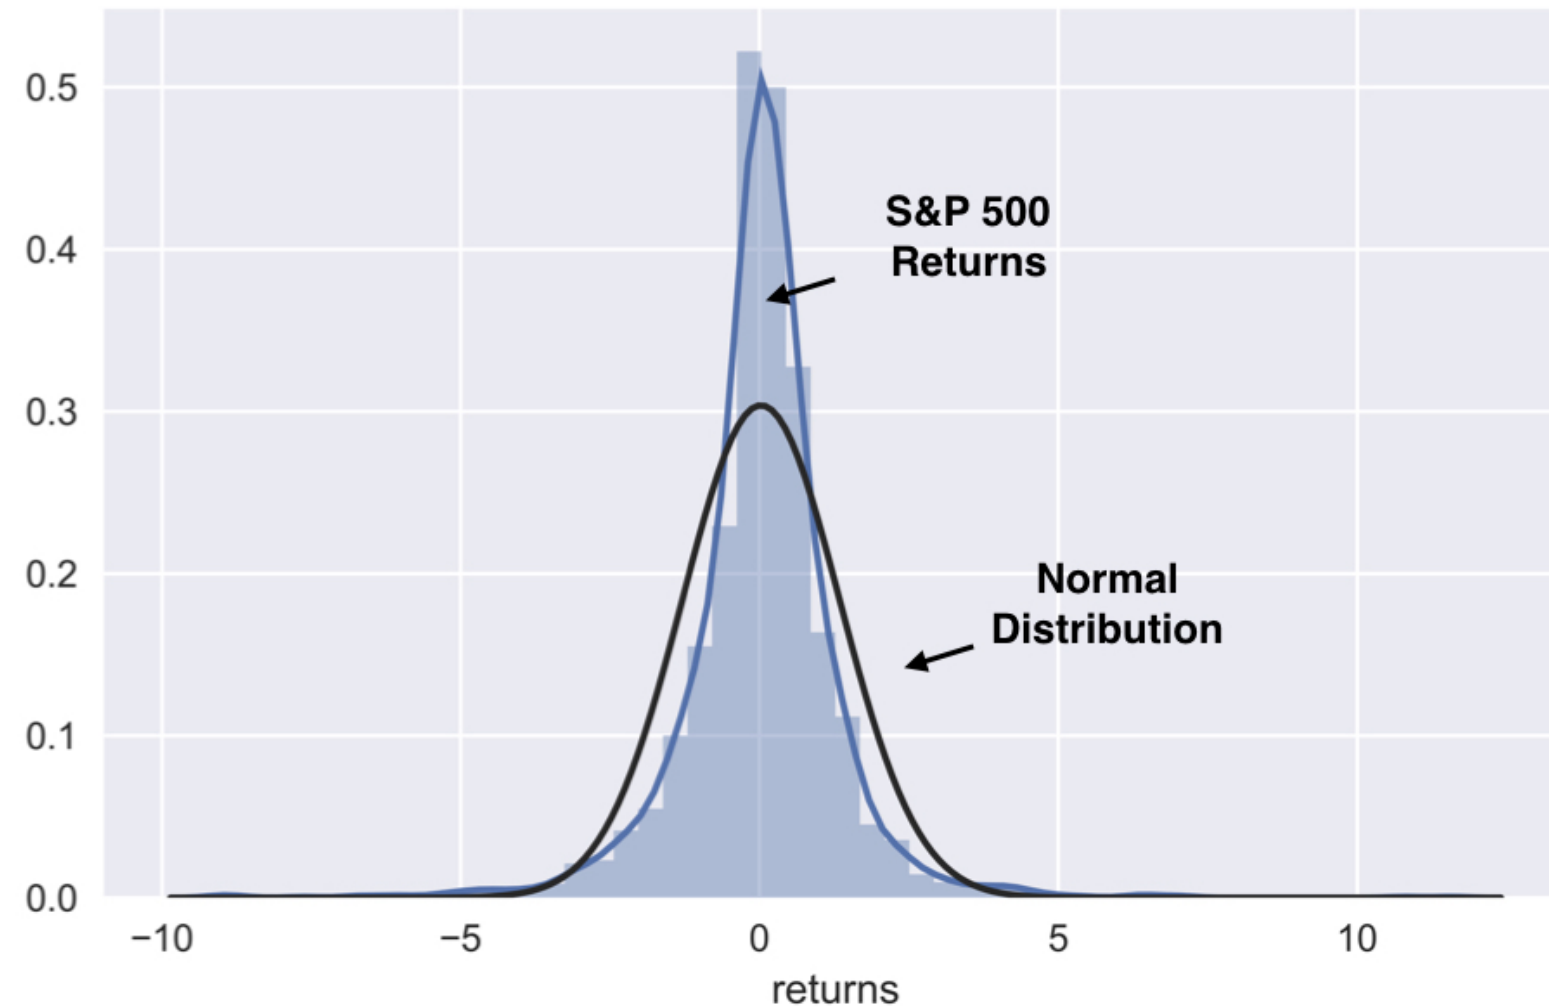
S&P 500 prices & returns

```
data = pd.read_csv('sp500.csv', parse_dates=['date'], index_col='date')  
data['returns'] = data.SP500.pct_change()  
data.plot(subplots=True)
```



S&P return distribution

```
sns.distplot(data.returns.dropna().mul(100), fit=norm)
```



Generate random S&P 500 returns

```
from numpy.random import choice
sample = data.returns.dropna()
n_obs = data.returns.count()
random_walk = choice(sample, size=n_obs)
random_walk = pd.Series(random_walk, index=sample.index)
random_walk.head()
```

```
DATE
2007-05-29    -0.008357
2007-05-30     0.003702
2007-05-31    -0.013990
2007-06-01     0.008096
2007-06-04     0.013120
```

Random S&P 500 prices (1)

```
start = data.SP500.first('D')
```


```
DATE
2007-05-25    1515.73
Name: SP500, dtype: float64
```

```
sp500_random = start.append(random_walk.add(1))
sp500_random.head()
```

```
DATE
2007-05-25    1515.730000
2007-05-29         0.998290
2007-05-30         0.995190
2007-05-31         0.997787
2007-06-01         0.983853
dtype: float64
```

Random S&P 500 prices (2)

```
data['SP500_random'] = sp500_random.cumprod()  
data[['SP500', 'SP500_random']].plot()
```



Let's practice!

MANIPULATING TIME SERIES DATA IN PYTHON

Relationships between time series: correlation

MANIPULATING TIME SERIES DATA IN PYTHON



Stefan Jansen

Founder & Lead Data Scientist at
Applied Artificial Intelligence

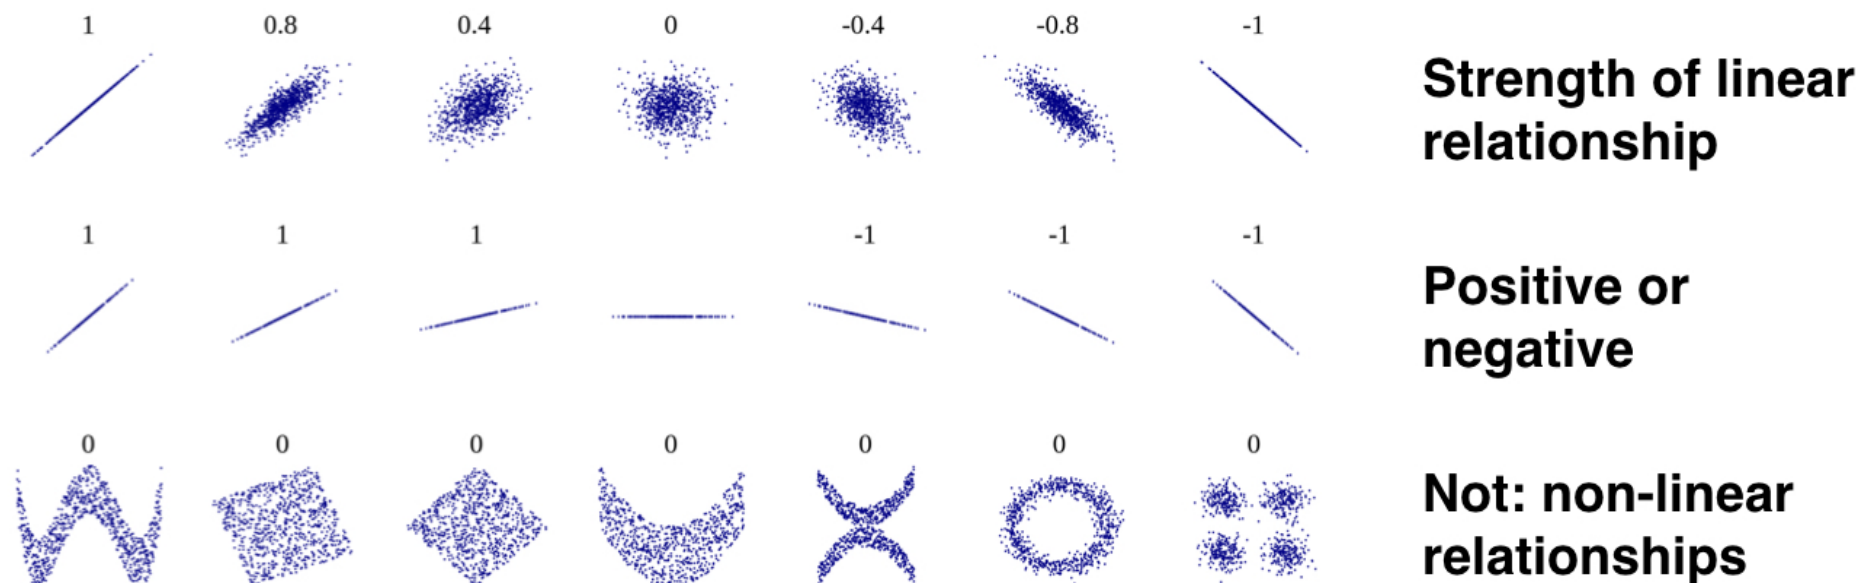
Correlation & relations between series

- So far, focus on characteristics of individual variables
- Now: characteristic of relations between variables
- Correlation: measures linear relationships
- Financial markets: important for prediction and risk management
- `pandas` & `seaborn` have tools to compute & visualize

Correlation & linear relationships

- Correlation coefficient: how similar is the pairwise movement of two variables around their averages?

- Varies between **-1** and **+1**
$$r = \frac{\sum_{i=1}^N (x_i - \bar{x})(y_i - \bar{y})}{s_x s_y}$$




Importing five price time series

```
data = pd.read_csv('assets.csv', parse_dates=['date'],  
                  index_col='date')  
  
data = data.dropna().info()
```

```
DatetimeIndex: 2469 entries, 2007-05-25 to 2017-05-22  
Data columns (total 5 columns):  
sp500      2469 non-null float64  
nasdaq     2469 non-null float64  
bonds      2469 non-null float64  
gold       2469 non-null float64  
oil        2469 non-null float64
```

Visualize pairwise linear relationships

```
daily_returns = data.pct_change()  
sns.jointplot(x='sp500', y='nasdaq', data=data_returns);
```




Calculate all correlations

```
correlations = returns.corr()  
correlations
```

	bonds	oil	gold	sp500	nasdaq
bonds	1.000000	-0.183755	0.003167	-0.300877	-0.306437
oil	-0.183755	1.000000	0.105930	0.335578	0.289590
gold	0.003167	0.105930	1.000000	-0.007786	-0.002544
sp500	-0.300877	0.335578	-0.007786	1.000000	0.959990
nasdaq	-0.306437	0.289590	-0.002544	0.959990	1.000000

Visualize all correlations

```
sns.heatmap(correlations, annot=True)
```



Let's practice!

MANIPULATING TIME SERIES DATA IN PYTHON