

# Streaming data case study

STREAMING DATA WITH AWS KINESIS AND LAMBDA



**Maksim Pecherskiy**  
Data Engineer

# This chapter


- Send incoming data to Firehose
- Store data
- Visualize data
- Set alerts in real-time
- Monitor the stream
- Meet a set of requirements




# Requirements

- Tweets must include #sandiego hashtag
- Tweets must come in real time
- Tweets must come enriched with sentiment
- Visualize last 15 minutes of data
- Notify manager if >3 negative tweets in 5 minute interval
- The stream should minimize data loss due to downtime
- Data must persist to be analyzed later

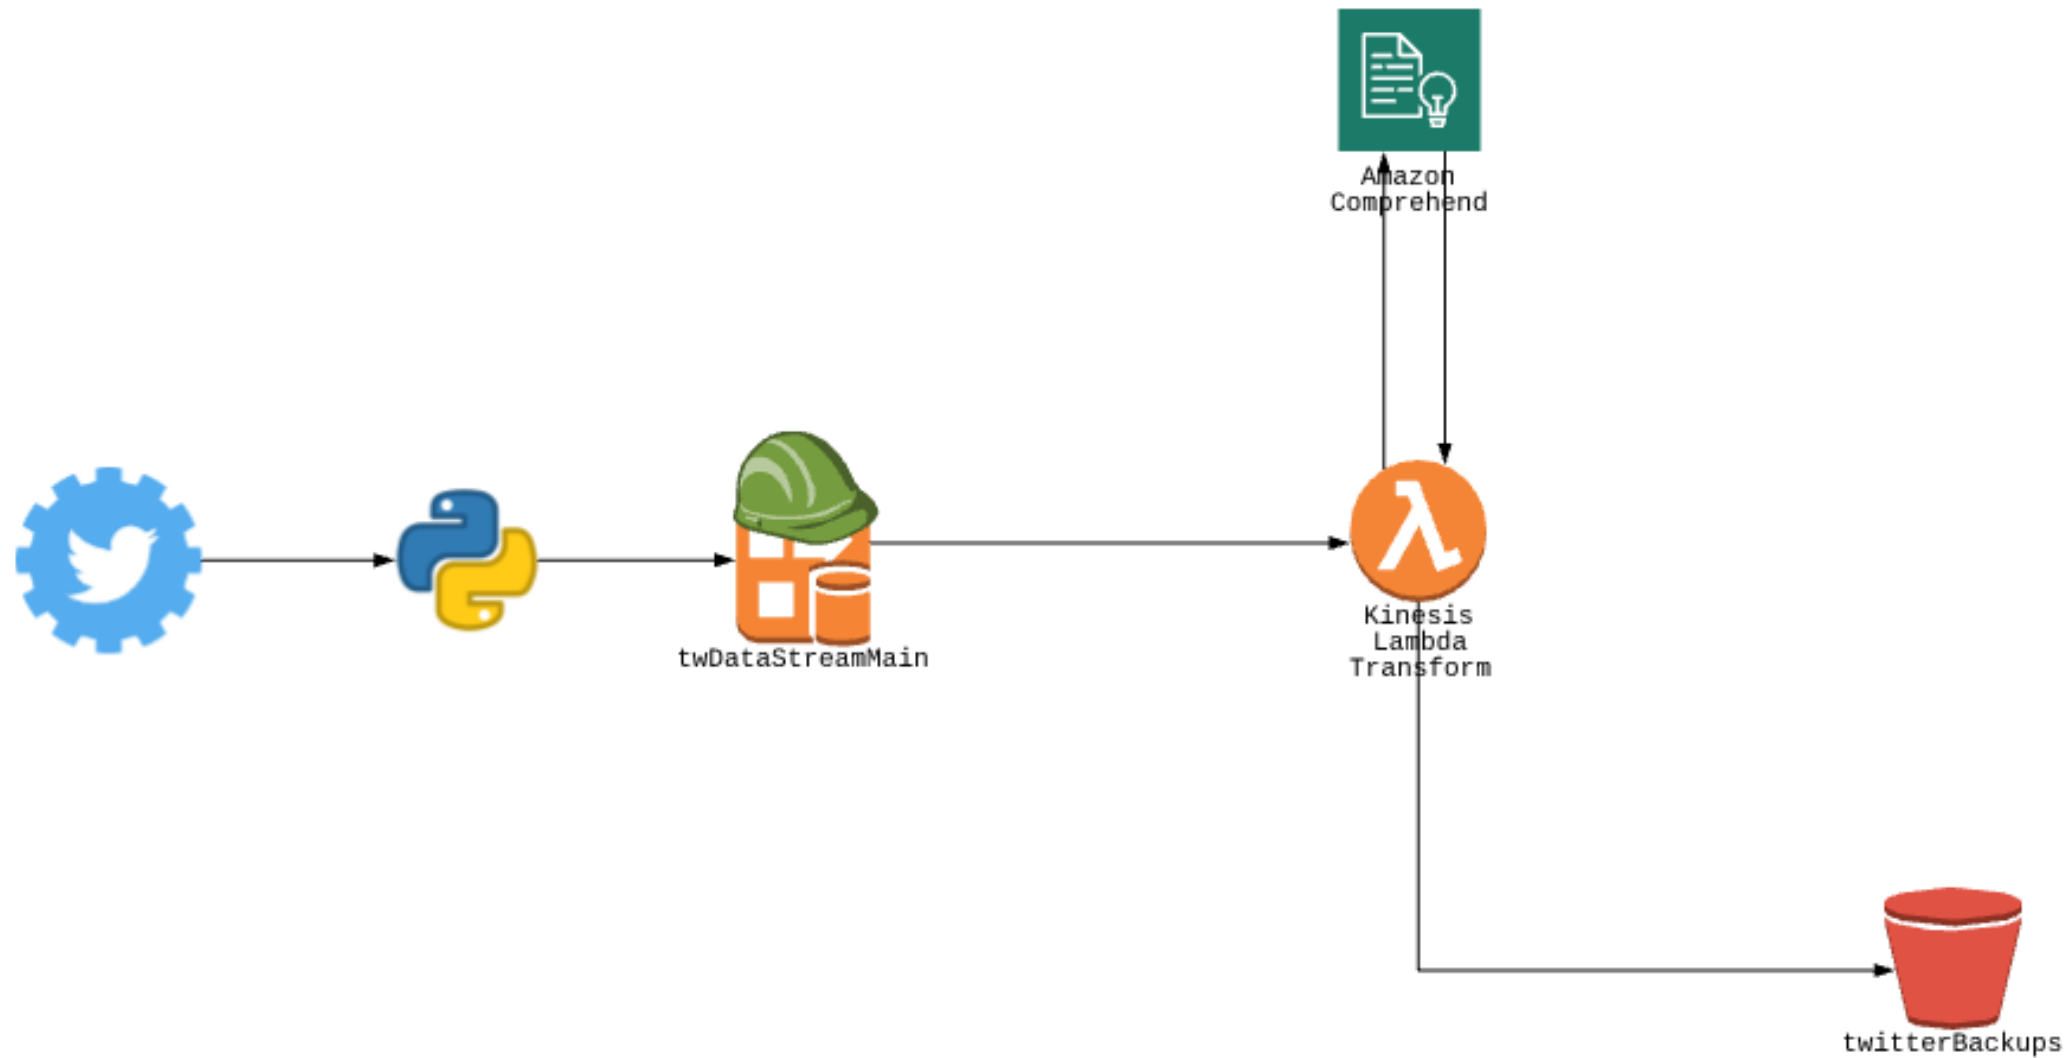
# Tweets come in real-time




# Enriched with sentiment




# Data must persist for later analysis



# Visualize last 15 minutes



Amazon  
Redshift



Amazon  
Elasticsearch  
Service



# Redshift vs Elasticsearch

## Redshift

- Designed for storing clean tables of data
- Schema is defined in database
- SQL for queries
- Works great with BI tools like Tableau

## Elasticsearch

- Schemaless - good for logs and text
- Schema is created during query
- Uses its own language for queries
- Has its own UI - Kibana

# Let's practice!

STREAMING DATA WITH AWS KINESIS AND LAMBDA

# Creating an Elasticsearch cluster

STREAMING DATA WITH AWS KINESIS AND LAMBDA



**Maksim Pecherskiy**  
Data Engineer

# Let's practice!

STREAMING DATA WITH AWS KINESIS AND LAMBDA

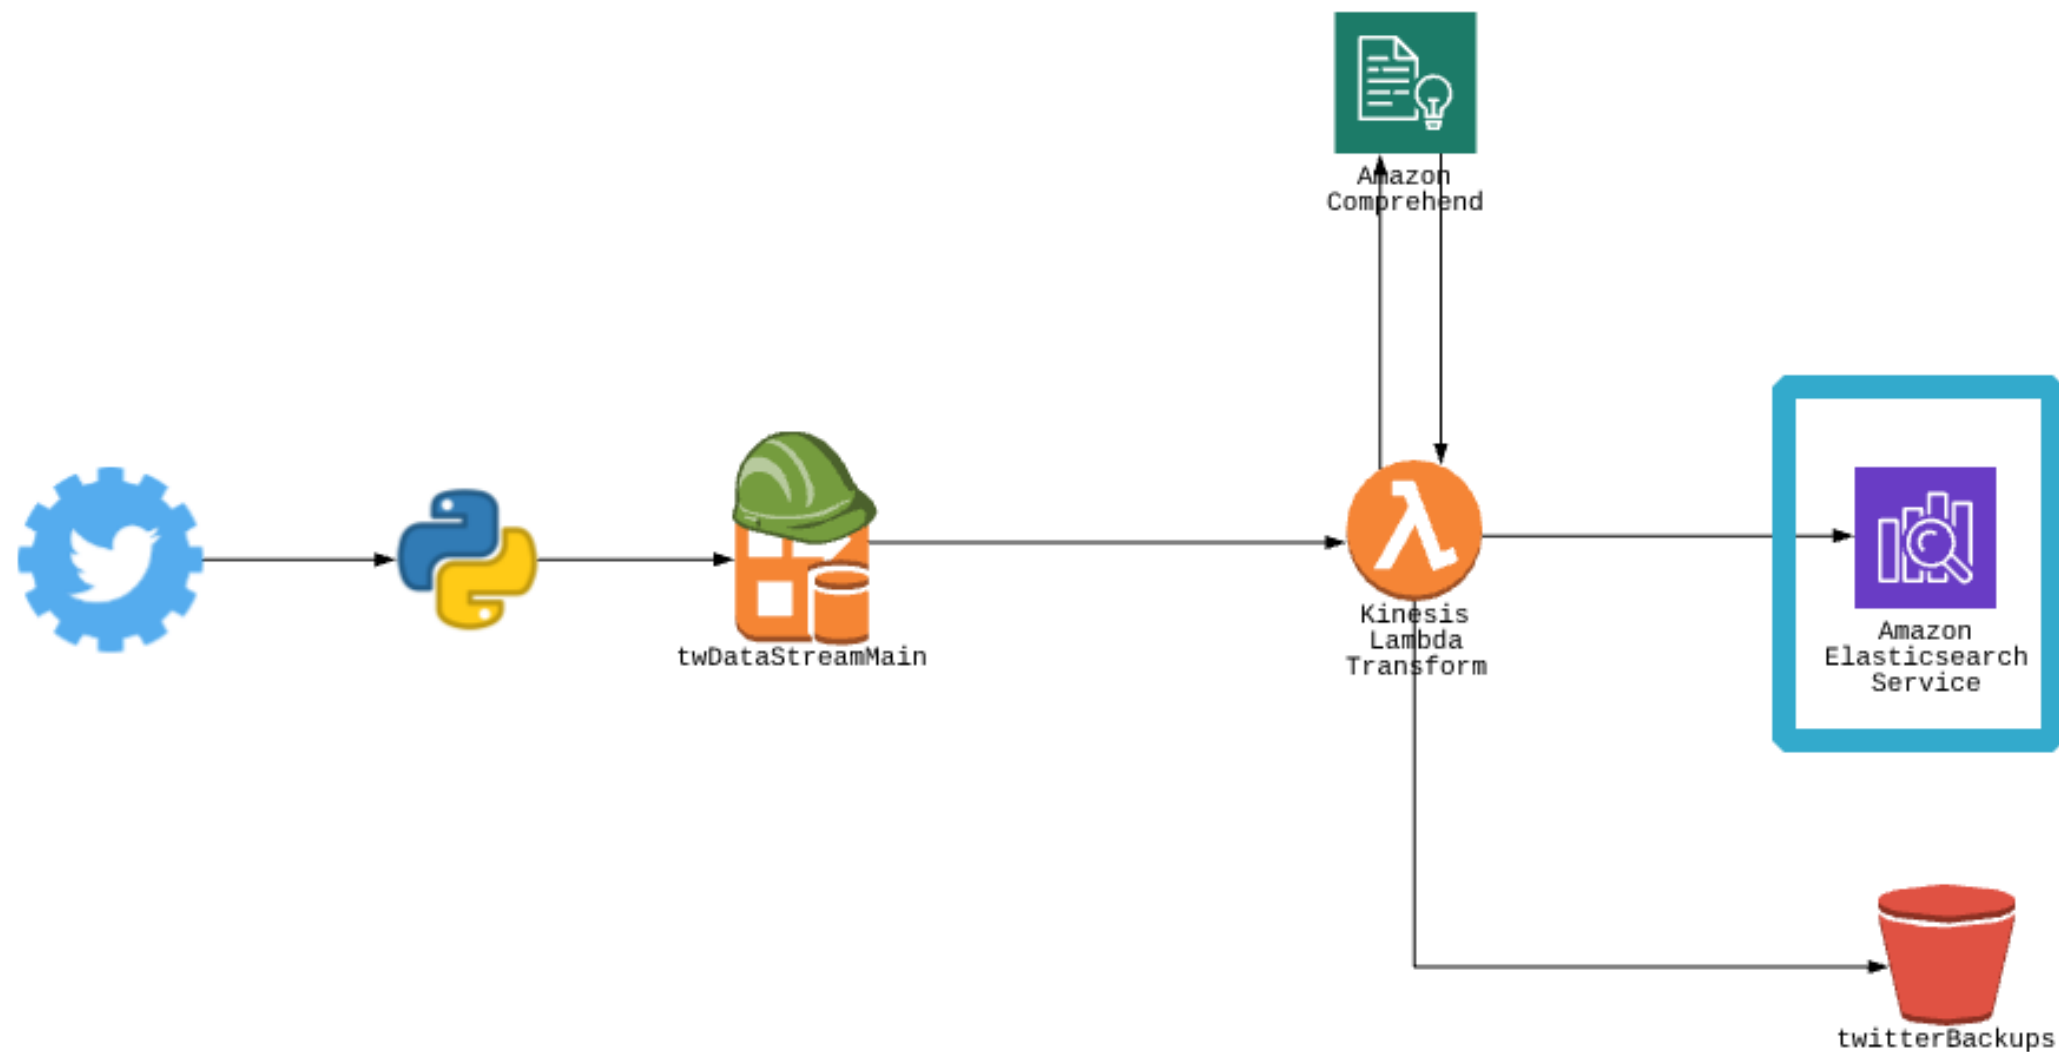
# Monitoring performance

STREAMING DATA WITH AWS KINESIS AND LAMBDA

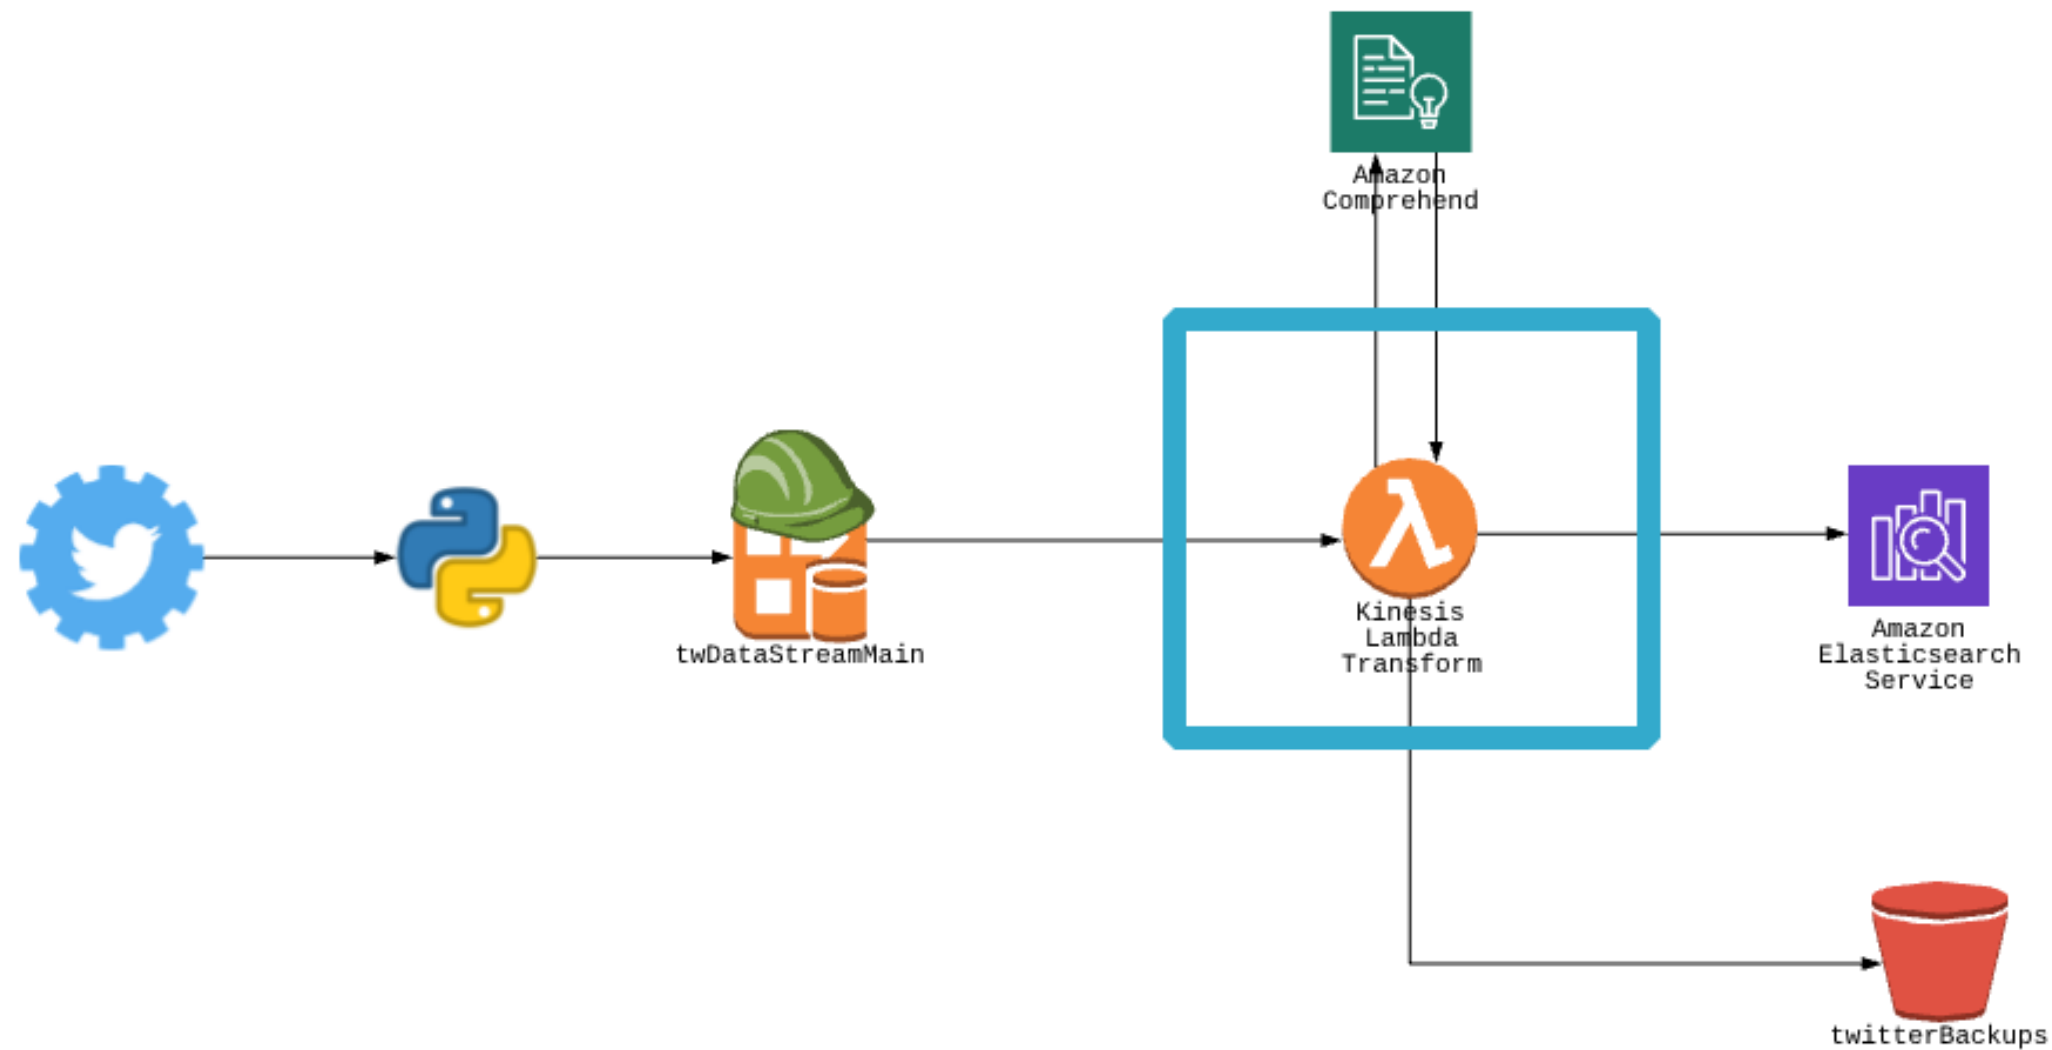


**Maksim Pecherskiy**  
Data Engineer

# Last lesson



# Lambda transform



# Lambda transform

```
def lambda_handler(event, context):  
    comprehend = boto3.client('comprehend',  
                               region_name='us-east-1',  
                               aws_access_key_id = AWS_KEY,  
                               aws_secret_access_key=AWS_SECRET)  
  
    output = []  
    for record in event['records']:  
        ...
```



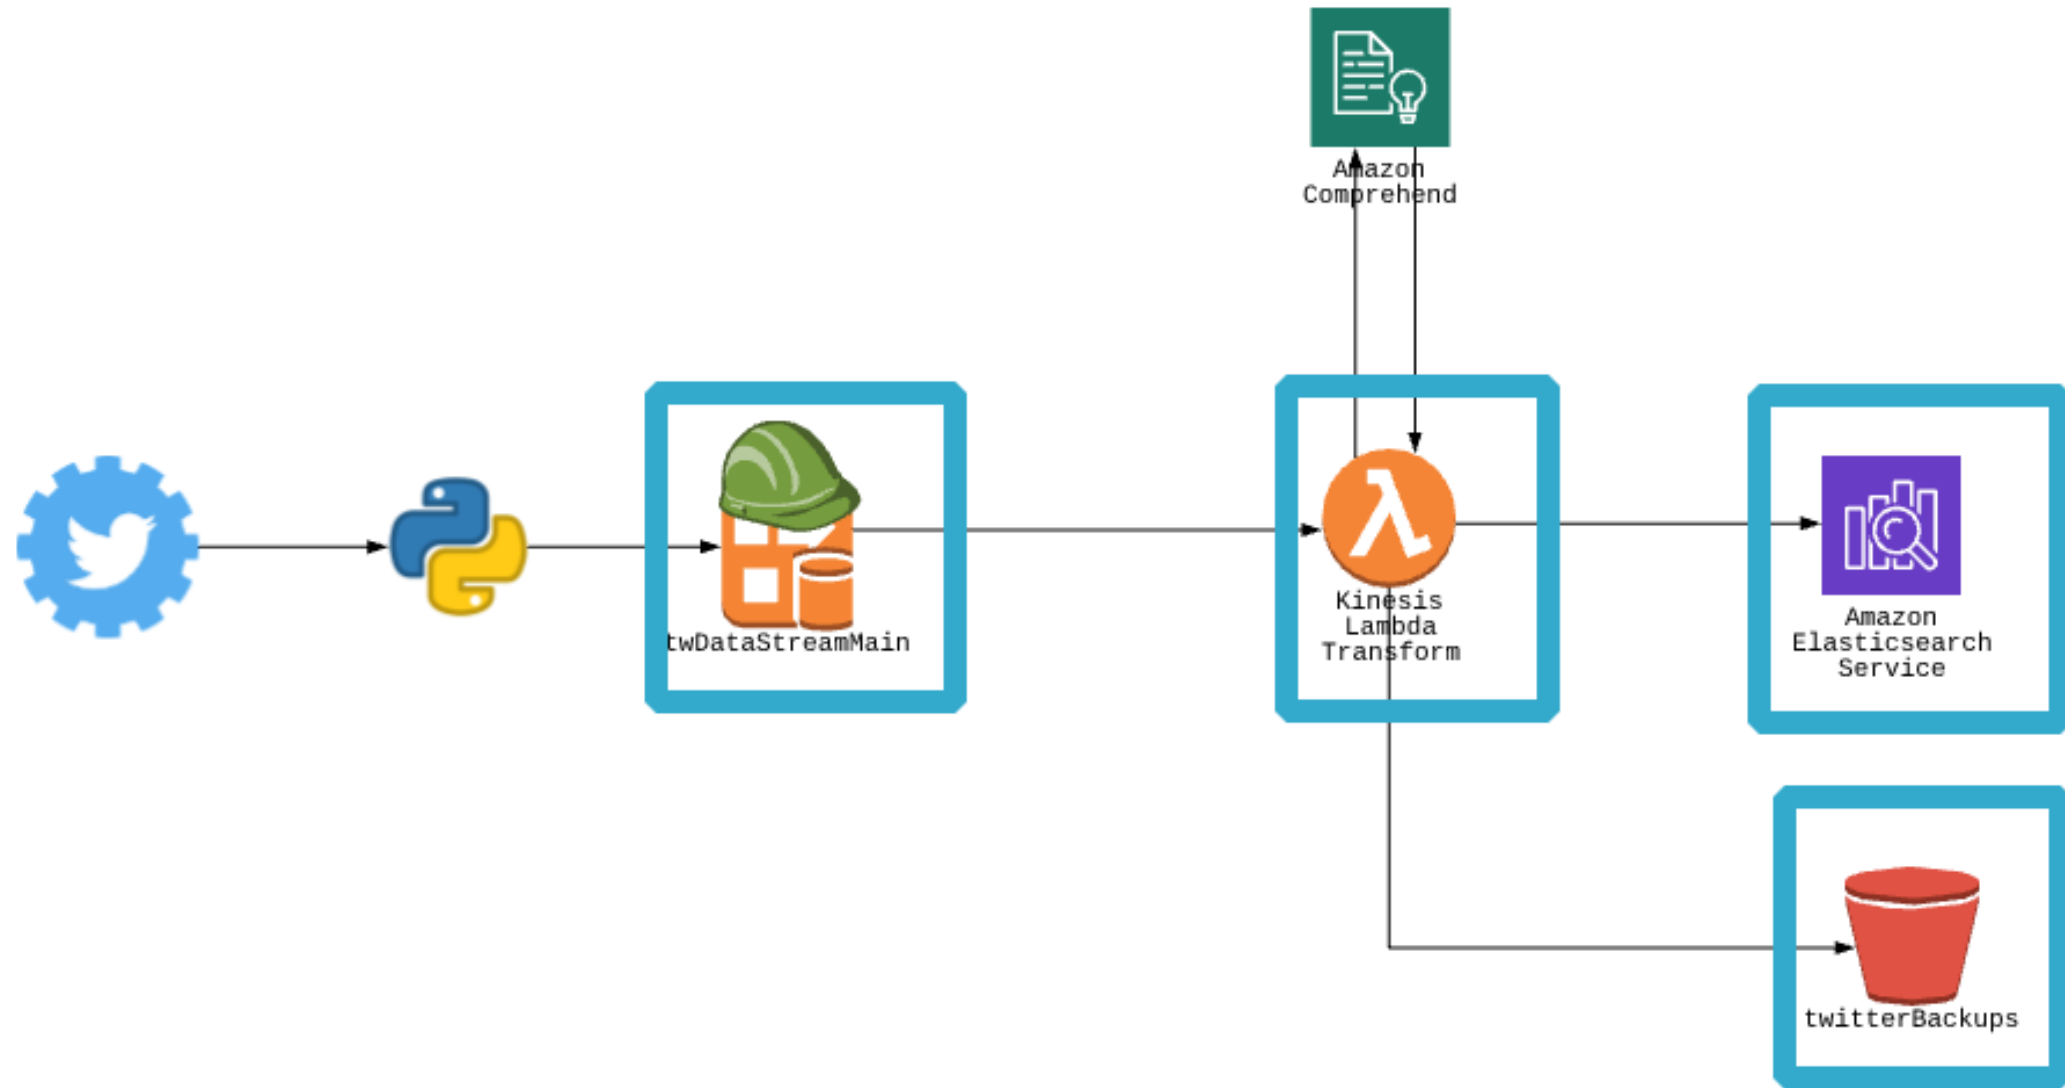
# Lambda transform

```
def lambda_handler(event, context):  
    ...  
    for record in event['records']:  
        dict_data = base64.b64decode(record['data']).decode('utf-8').strip()  
        dict_data = json.loads(dict_data)  
        sentiment_all = comprehend.detect_sentiment(  
            Text=dict_data['text'],  
            LanguageCode=dict_data['lang'])  
        dict_data['sentiment'] = sentiment_all['Sentiment']  
    ...
```

# Lambda transform

```
def lambda_handler(event, context):  
    ...  
    for record in event['records']:  
        ...  
        output_record = {  
            'recordId': record['recordId'],  
            'result': 'Ok',  
            'data': base64.b64encode(json.dumps(dict_data).encode('utf-8'))  
        }  
        output.append(output_record)  
    return {'records': output}
```

# Wiring it up





# Update firehoseDeliveryRole

[Roles](#) > [firehoseDeliveryRole](#)

## Summary

Delete role

Role ARN	arn:aws:iam::458913182630:role/firehoseDeliveryRole 
Role description	<a href="#">Edit</a>
Instance Profile ARNs	
Path	/
Creation time	2020-04-13 08:54 EDT
Last activity	2020-07-26 16:38 EDT (7 days ago)
Maximum session duration	1 hour <a href="#">Edit</a>

Permissions

Trust relationships

Tags

Access Advisor

Revoke sessions

▼ Permissions policies (3 policies applied)

Attach policies

[+ Add inline policy](#)


Policy name ▼	Policy type ▼	
▶  <a href="#">AWSLambdaFullAccess</a>	AWS managed policy	✕
▶  <a href="#">AmazonS3FullAccess</a>	AWS managed policy	✕
▶  <a href="#">AmazonESFullAccess</a>	AWS managed policy	✕

# Create delivery stream

## Amazon Elasticsearch Service destination

---

### Domain

You can select a domain that resides within a VPC or one that uses a public endpoint. If your domain uses a public endpoint, you don't need to configure this delivery stream. [Learn more](#) 

tw-data-domain-1



Create new 

View **tw-data-domain-1** in Amazon Elasticsearch Service 

### Index

sd\_tweets\_tr

A new index will be created if the the specified index name does not exist.

# Create delivery stream

## S3 backup

Backup mode

All records

S3 bucket

tw-backups-33



Create new

[View tw-backups-33 in S3 console](#)

Backup S3 bucket prefix - *optional*

comicon-tweets-bk/

Buffer size

1

MiB

Enter a buffer size between 1 - 128 MiB.

Buffer interval

60

seconds

Enter a buffer interval between 60 - 900 seconds.

S3 Compression

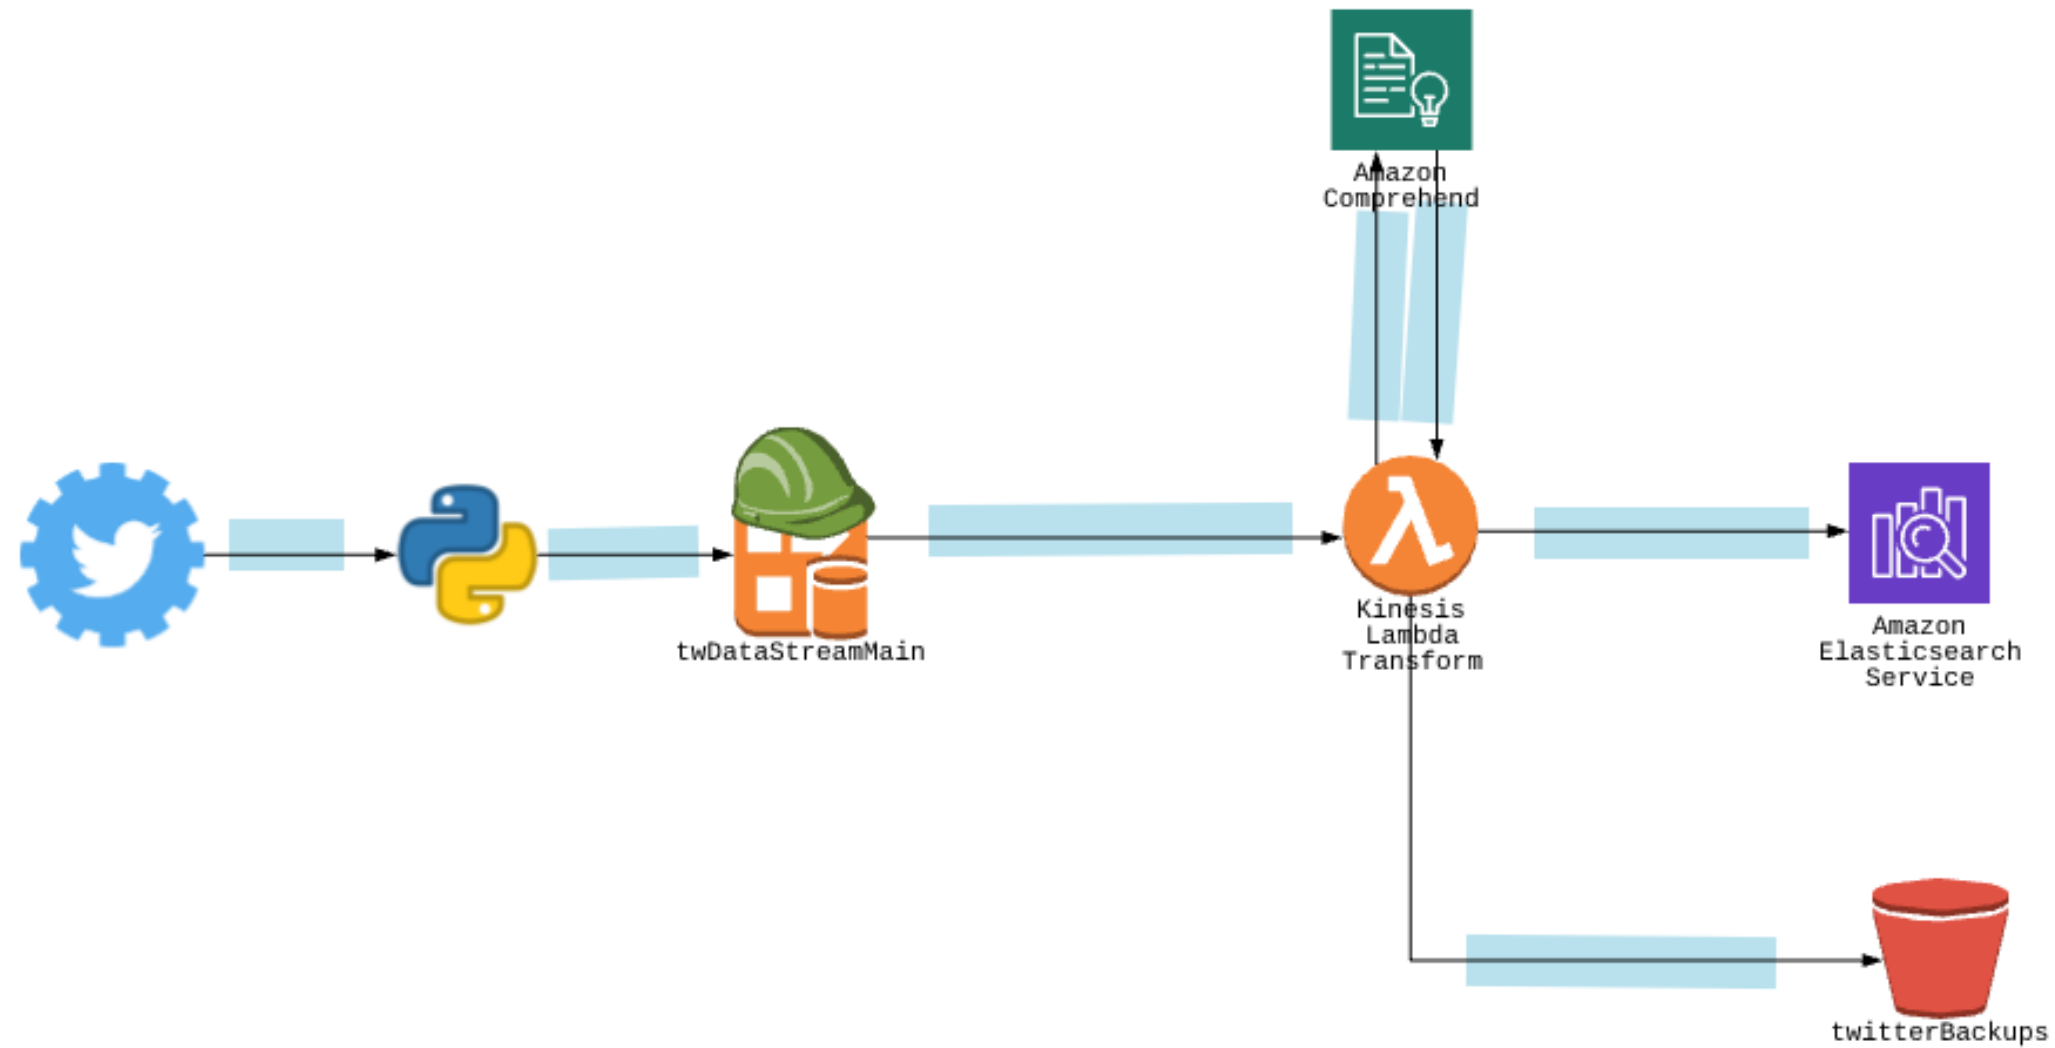
Standard

# Cloudwatch

The stream should minimize data loss due to downtime.

- Logs (raw data)
- Metrics (measures of various activities of the service)
- Alarms (notifications when a metric is out of a specified range)
- Dashboards (metrics visualization)

# Failure points





# Let's practice!

STREAMING DATA WITH AWS KINESIS AND LAMBDA

# Cloudwatch dashboards and alarms

STREAMING DATA WITH AWS KINESIS AND LAMBDA



**Maksim Pecherskiy**  
Data Engineer

# Let's practice!

STREAMING DATA WITH AWS KINESIS AND LAMBDA

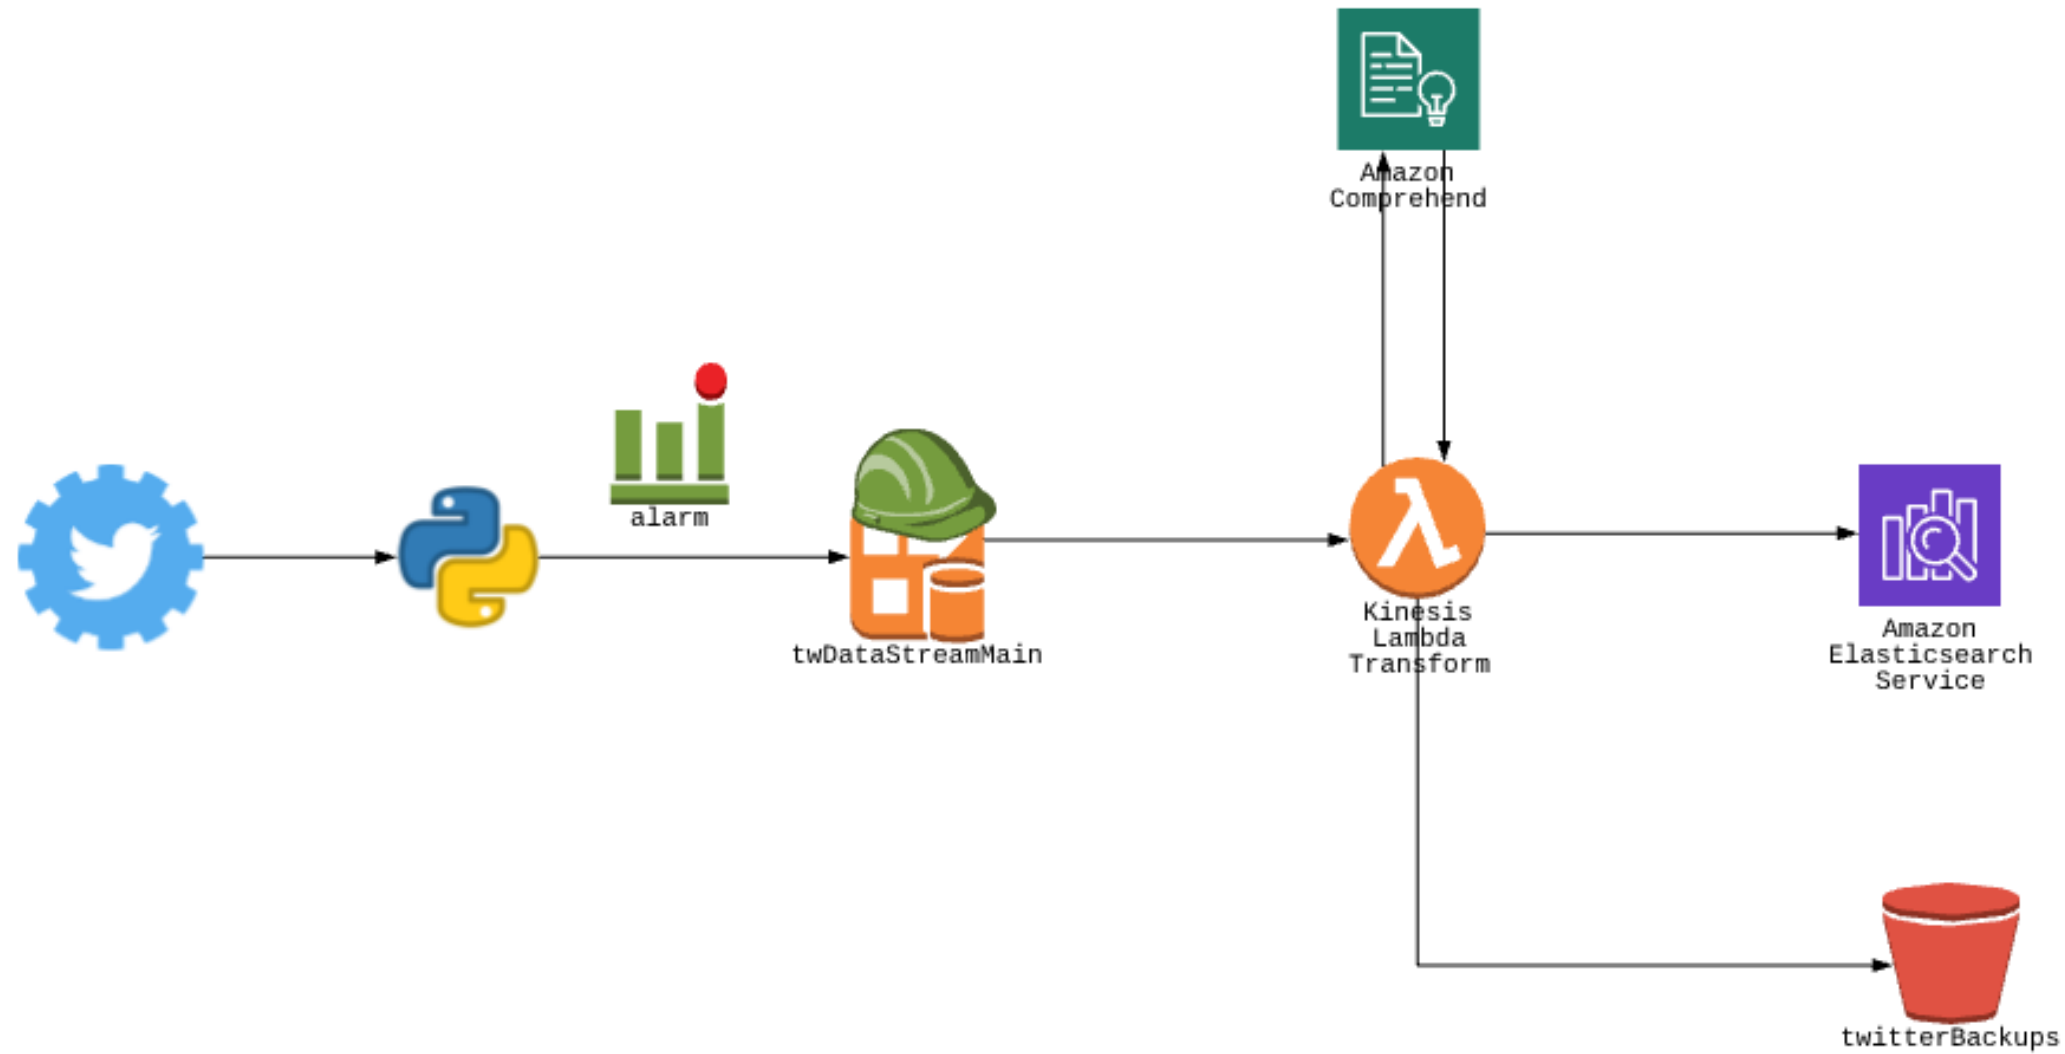
# Visualizing streaming data

STREAMING DATA WITH AWS KINESIS AND LAMBDA

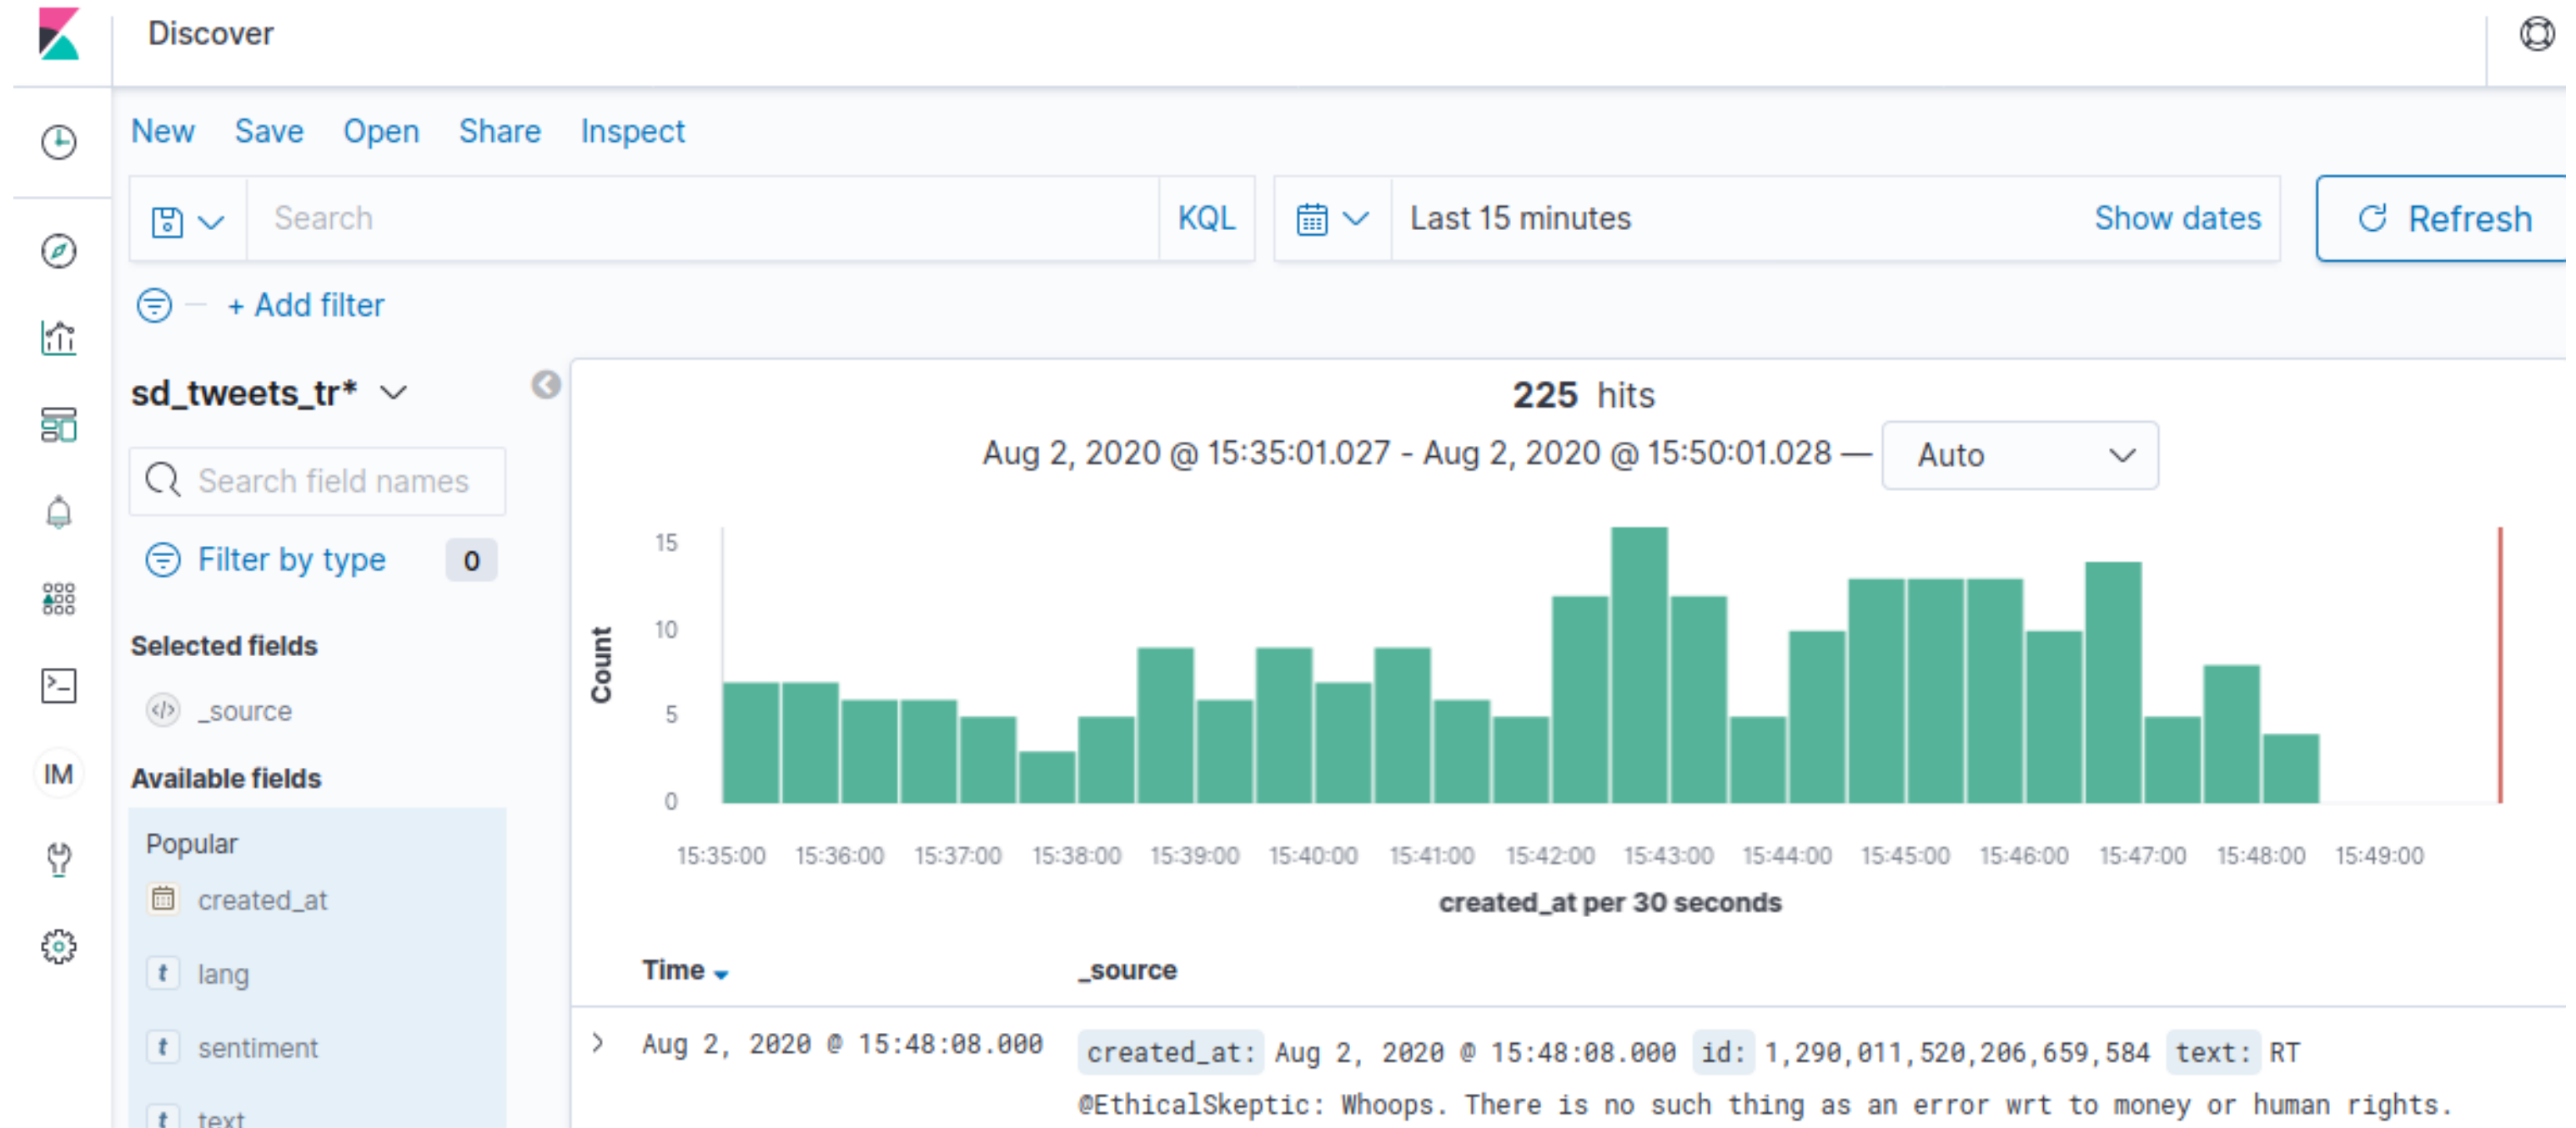


**Maksim Pecherskiy**  
Data Engineer

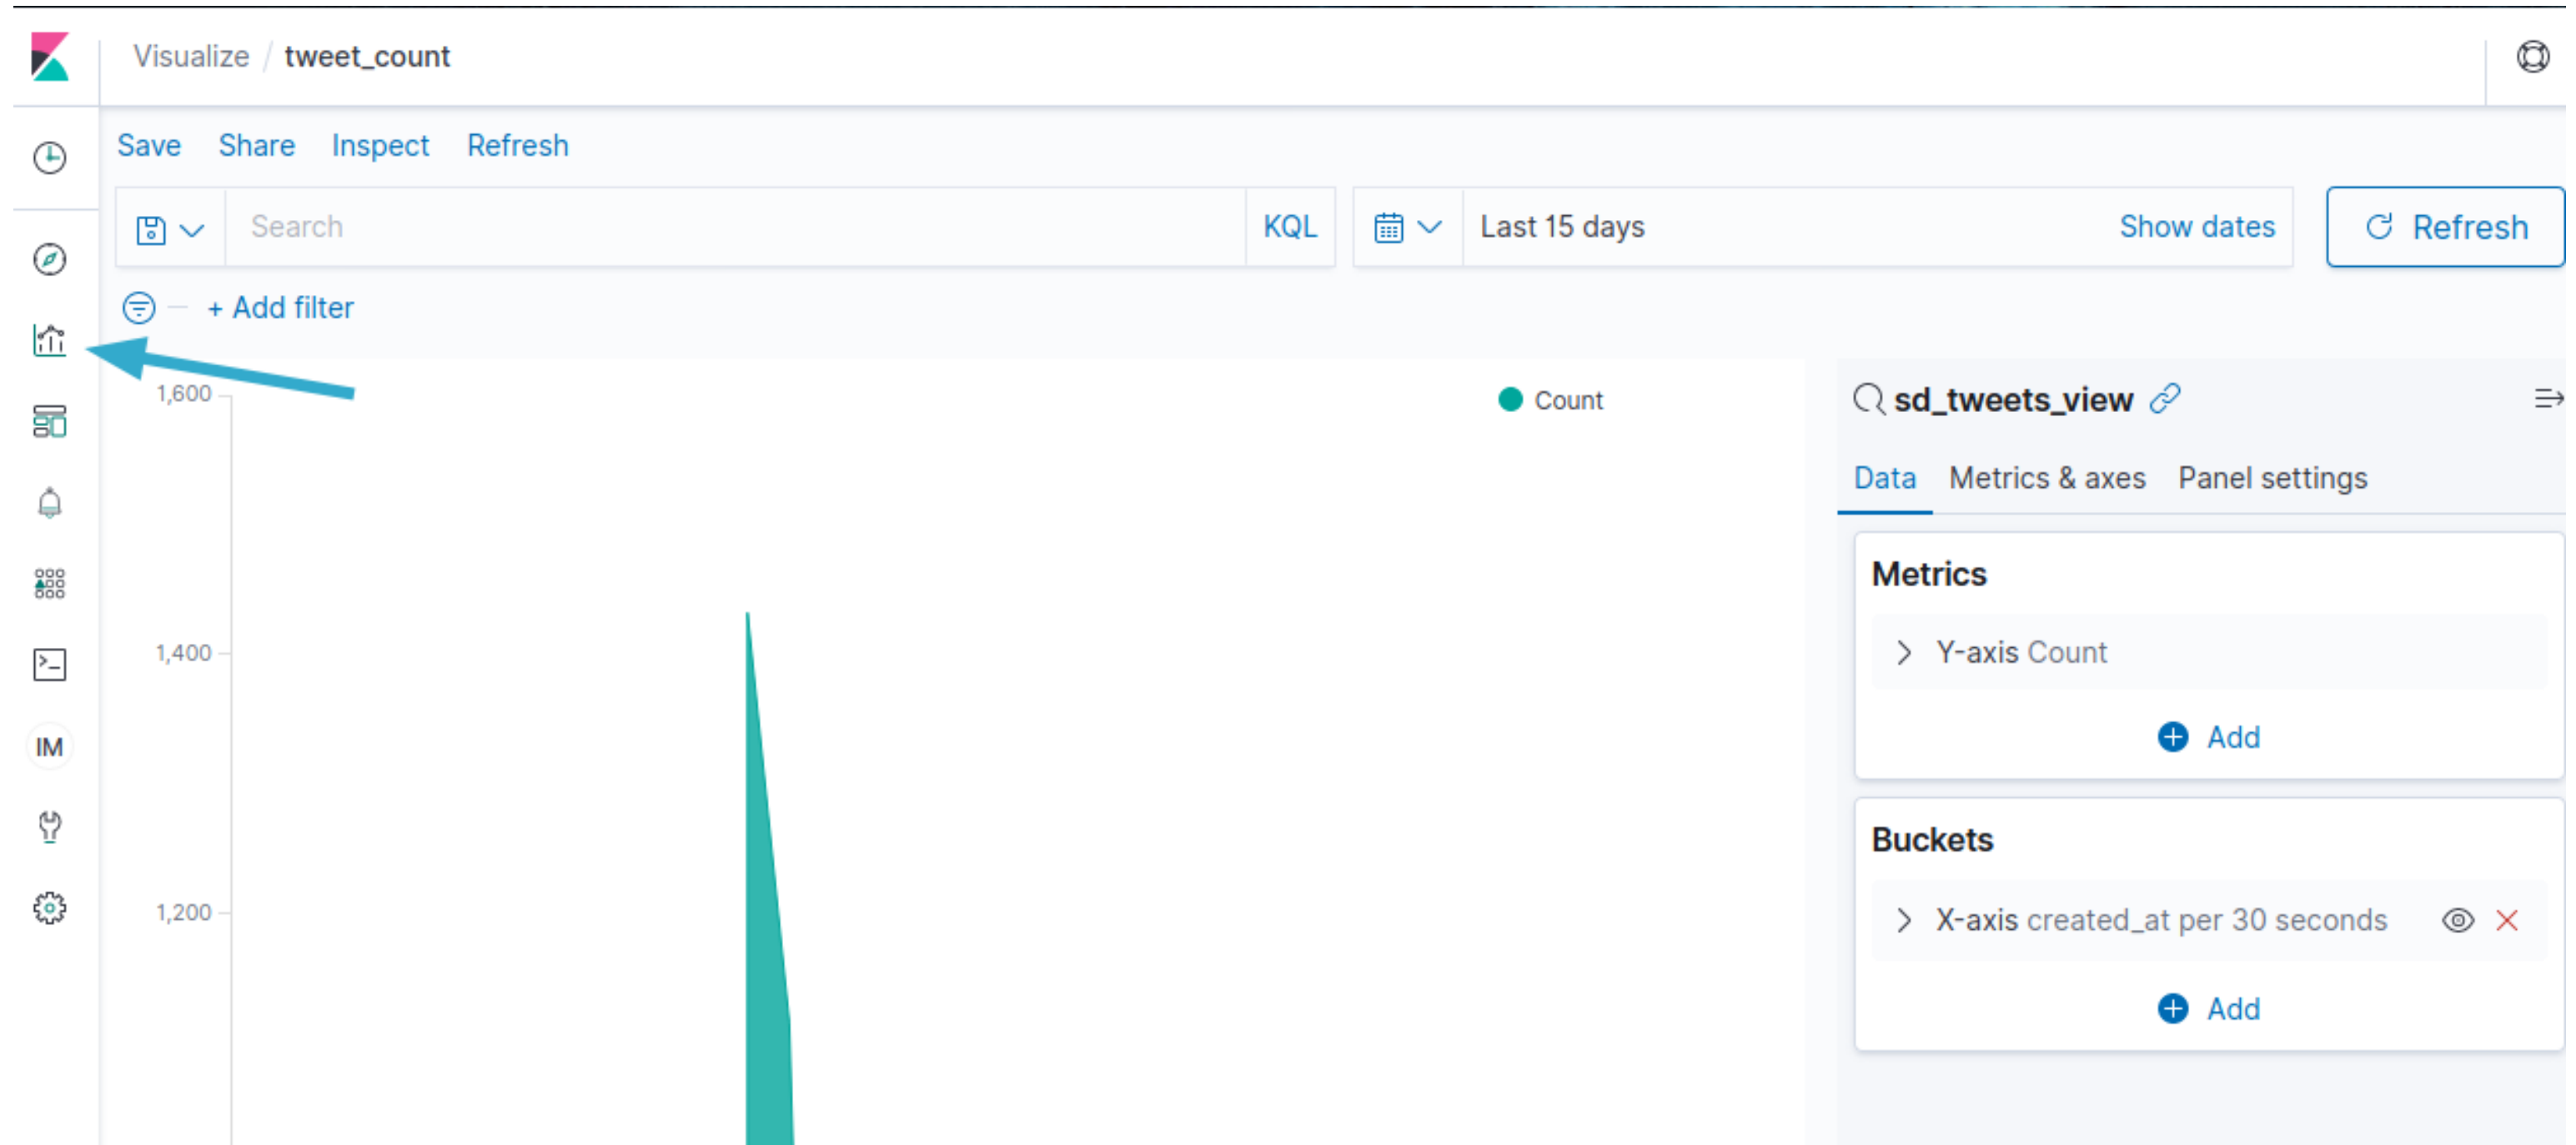
# Our pipeline so far



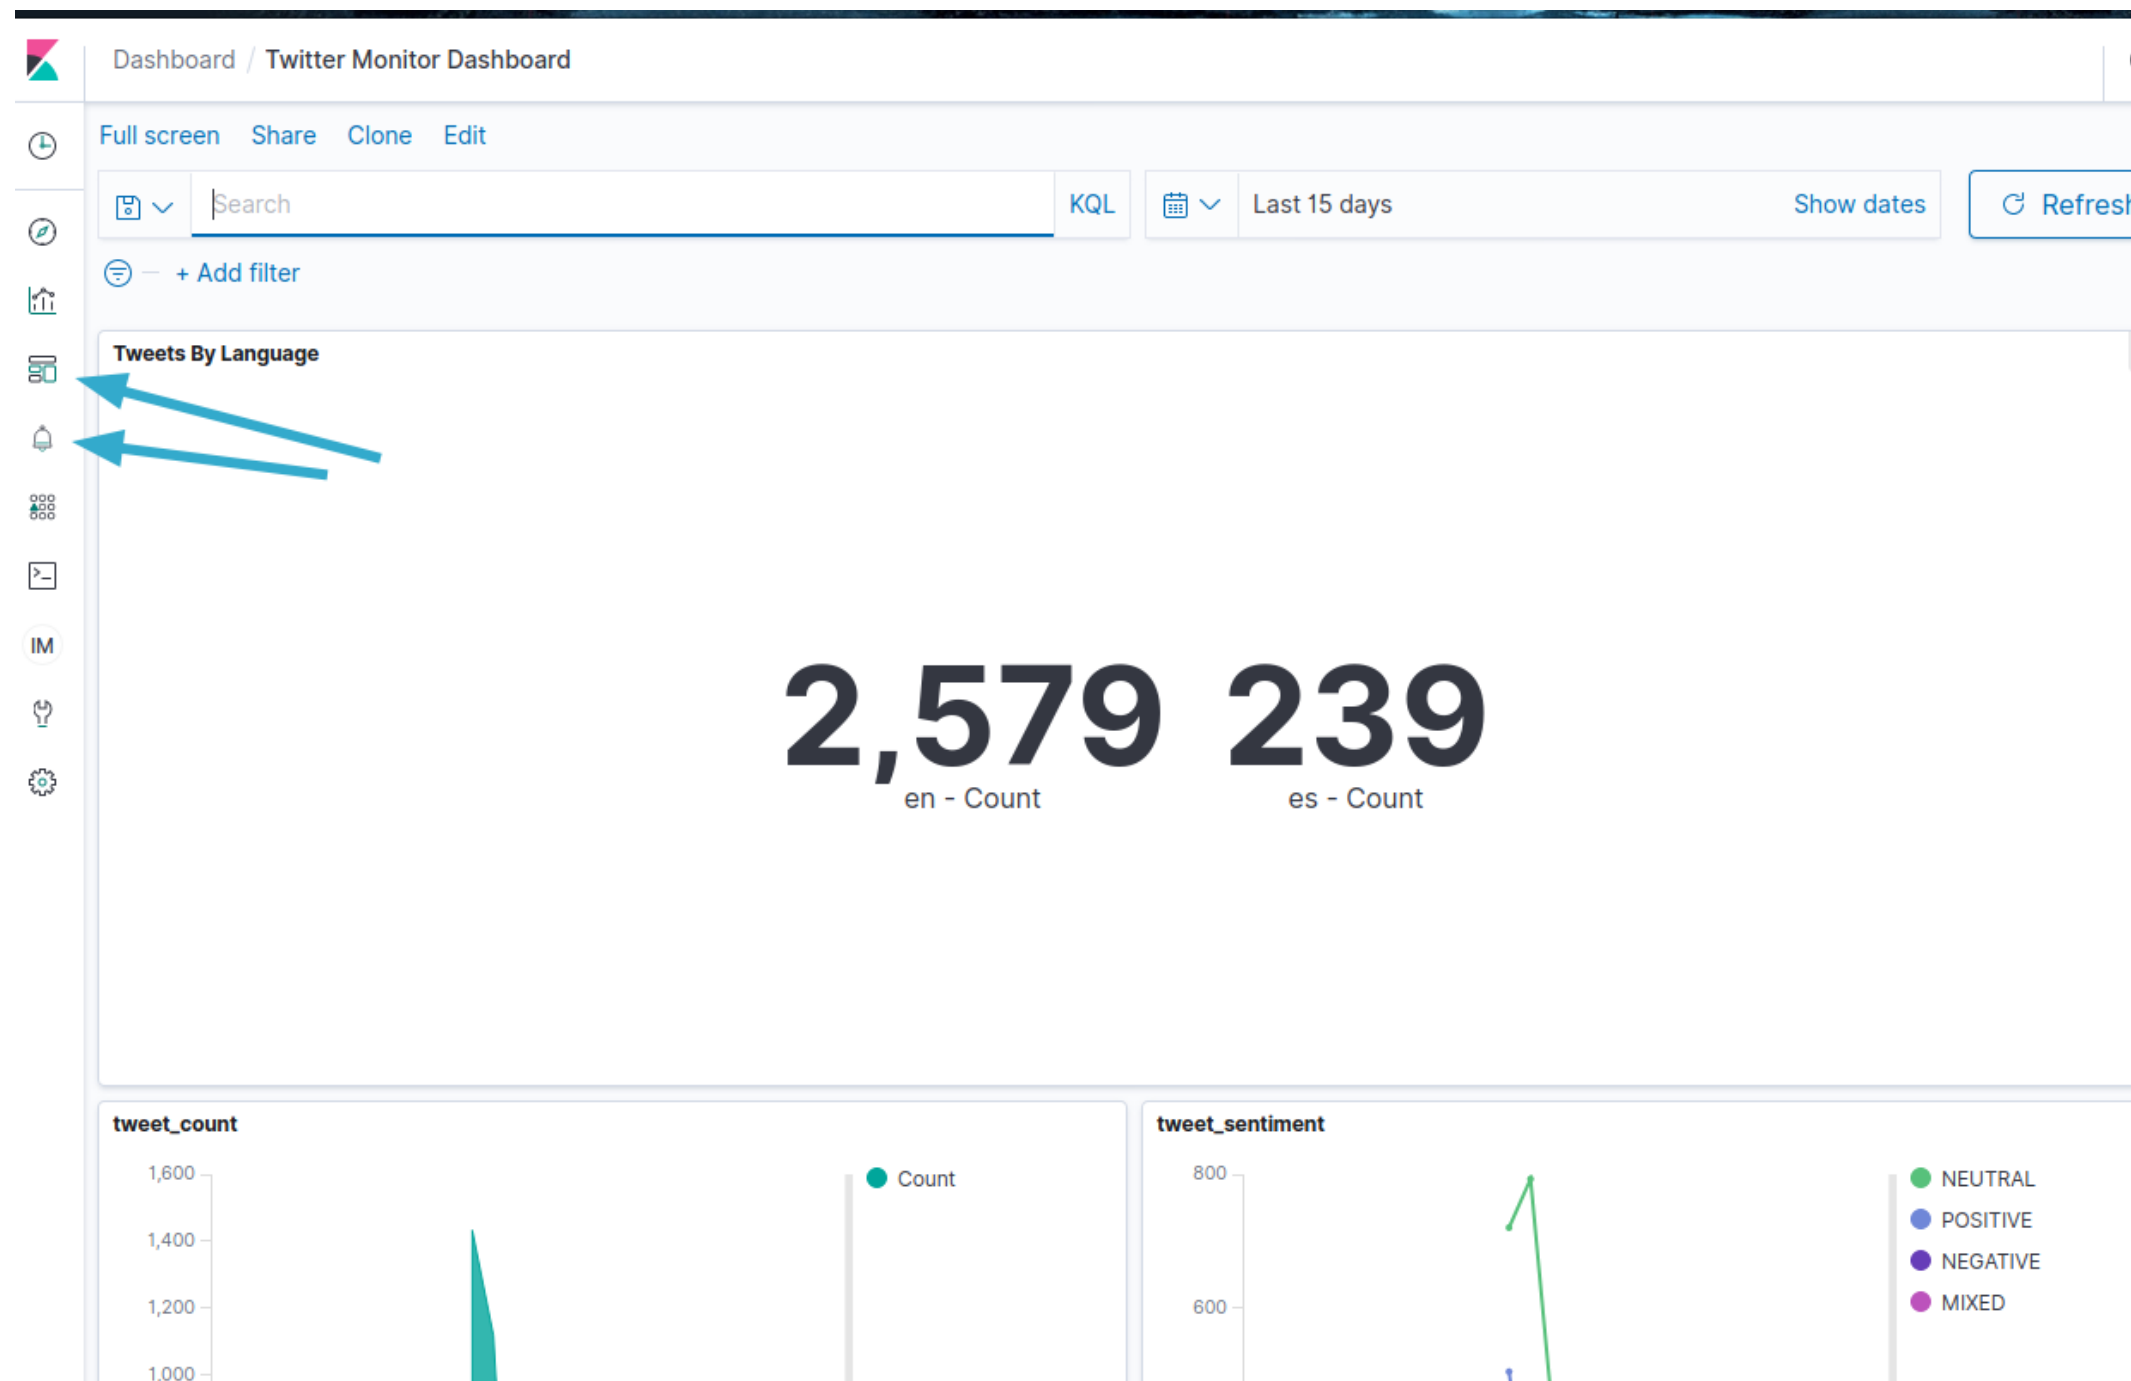
# Kibana Discover view



# Kibana visualizations



# Dashboards and alerts





# Elasticsearch vs CloudWatch

## CloudWatch

- AWS centric
- Can accept custom data, but not primary use case
- Great for working with logs

## Elasticsearch

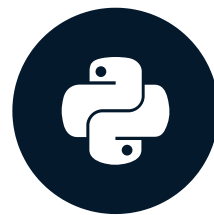
- General, open source tool
- Accepts custom data, including logs
- Better viz than CloudWatch
- Robust plugin ecosystem

# Let's practice!

STREAMING DATA WITH AWS KINESIS AND LAMBDA

# Working with ElasticSearch using Kibana

STREAMING DATA WITH AWS KINESIS AND LAMBDA



**Maksim Pecherskiy**  
Data Engineer

# Let's practice!

STREAMING DATA WITH AWS KINESIS AND LAMBDA