Defining neural networks with Keras
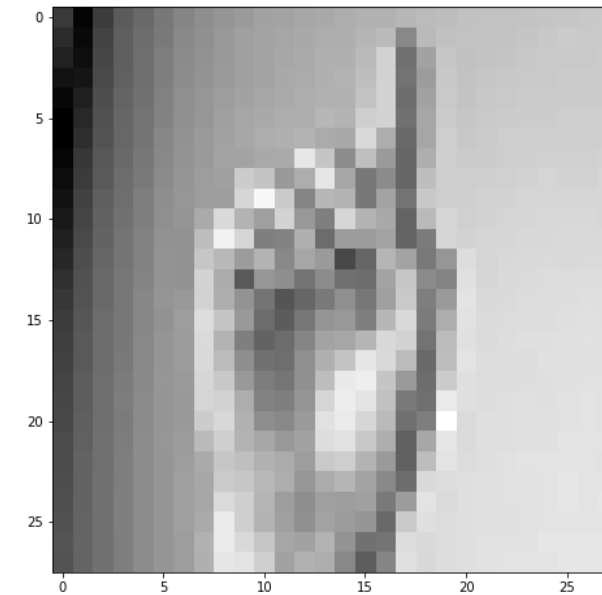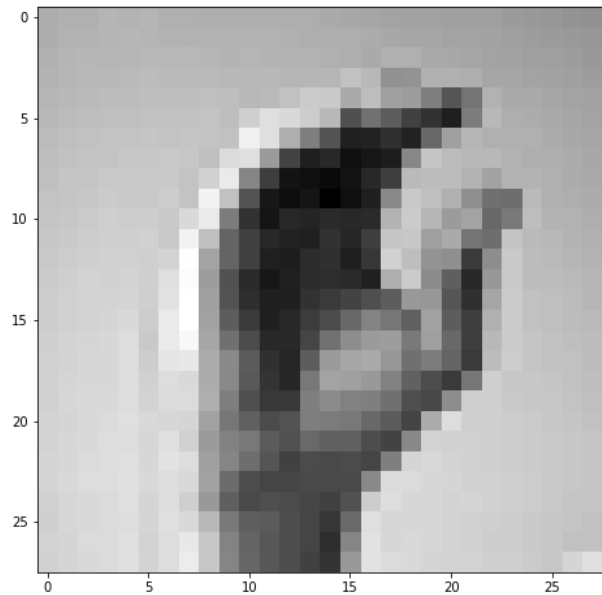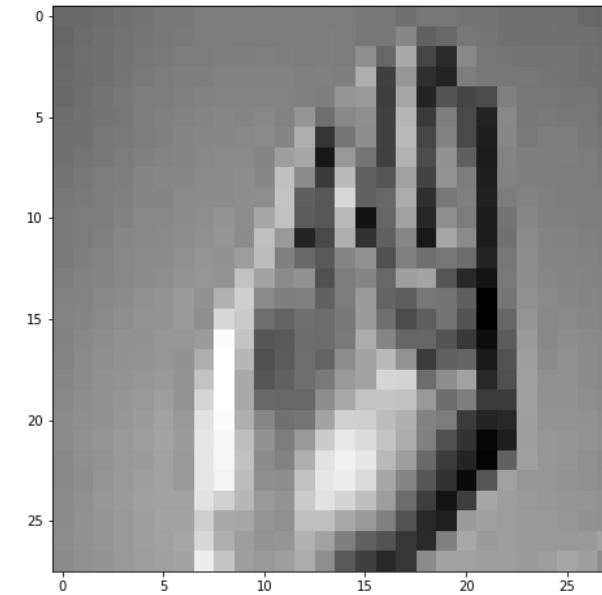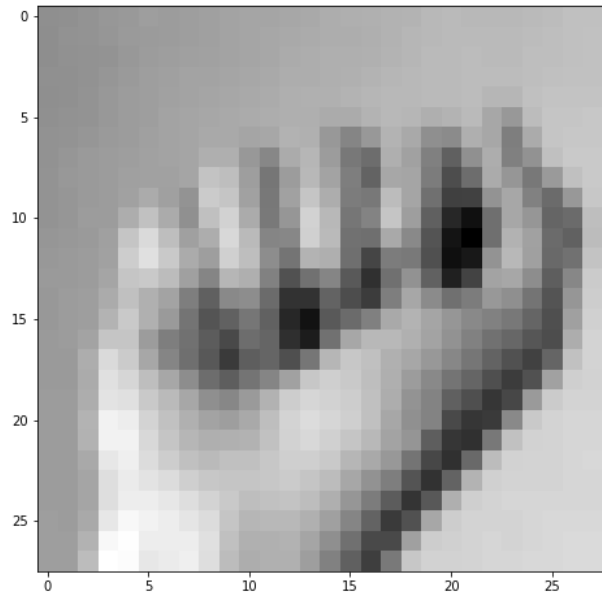
INTRODUCTION TO TENSORFLOW IN PYTHON

Isaiah Hull

Economist

datacamp

Classifying sign language letters

The sequential API
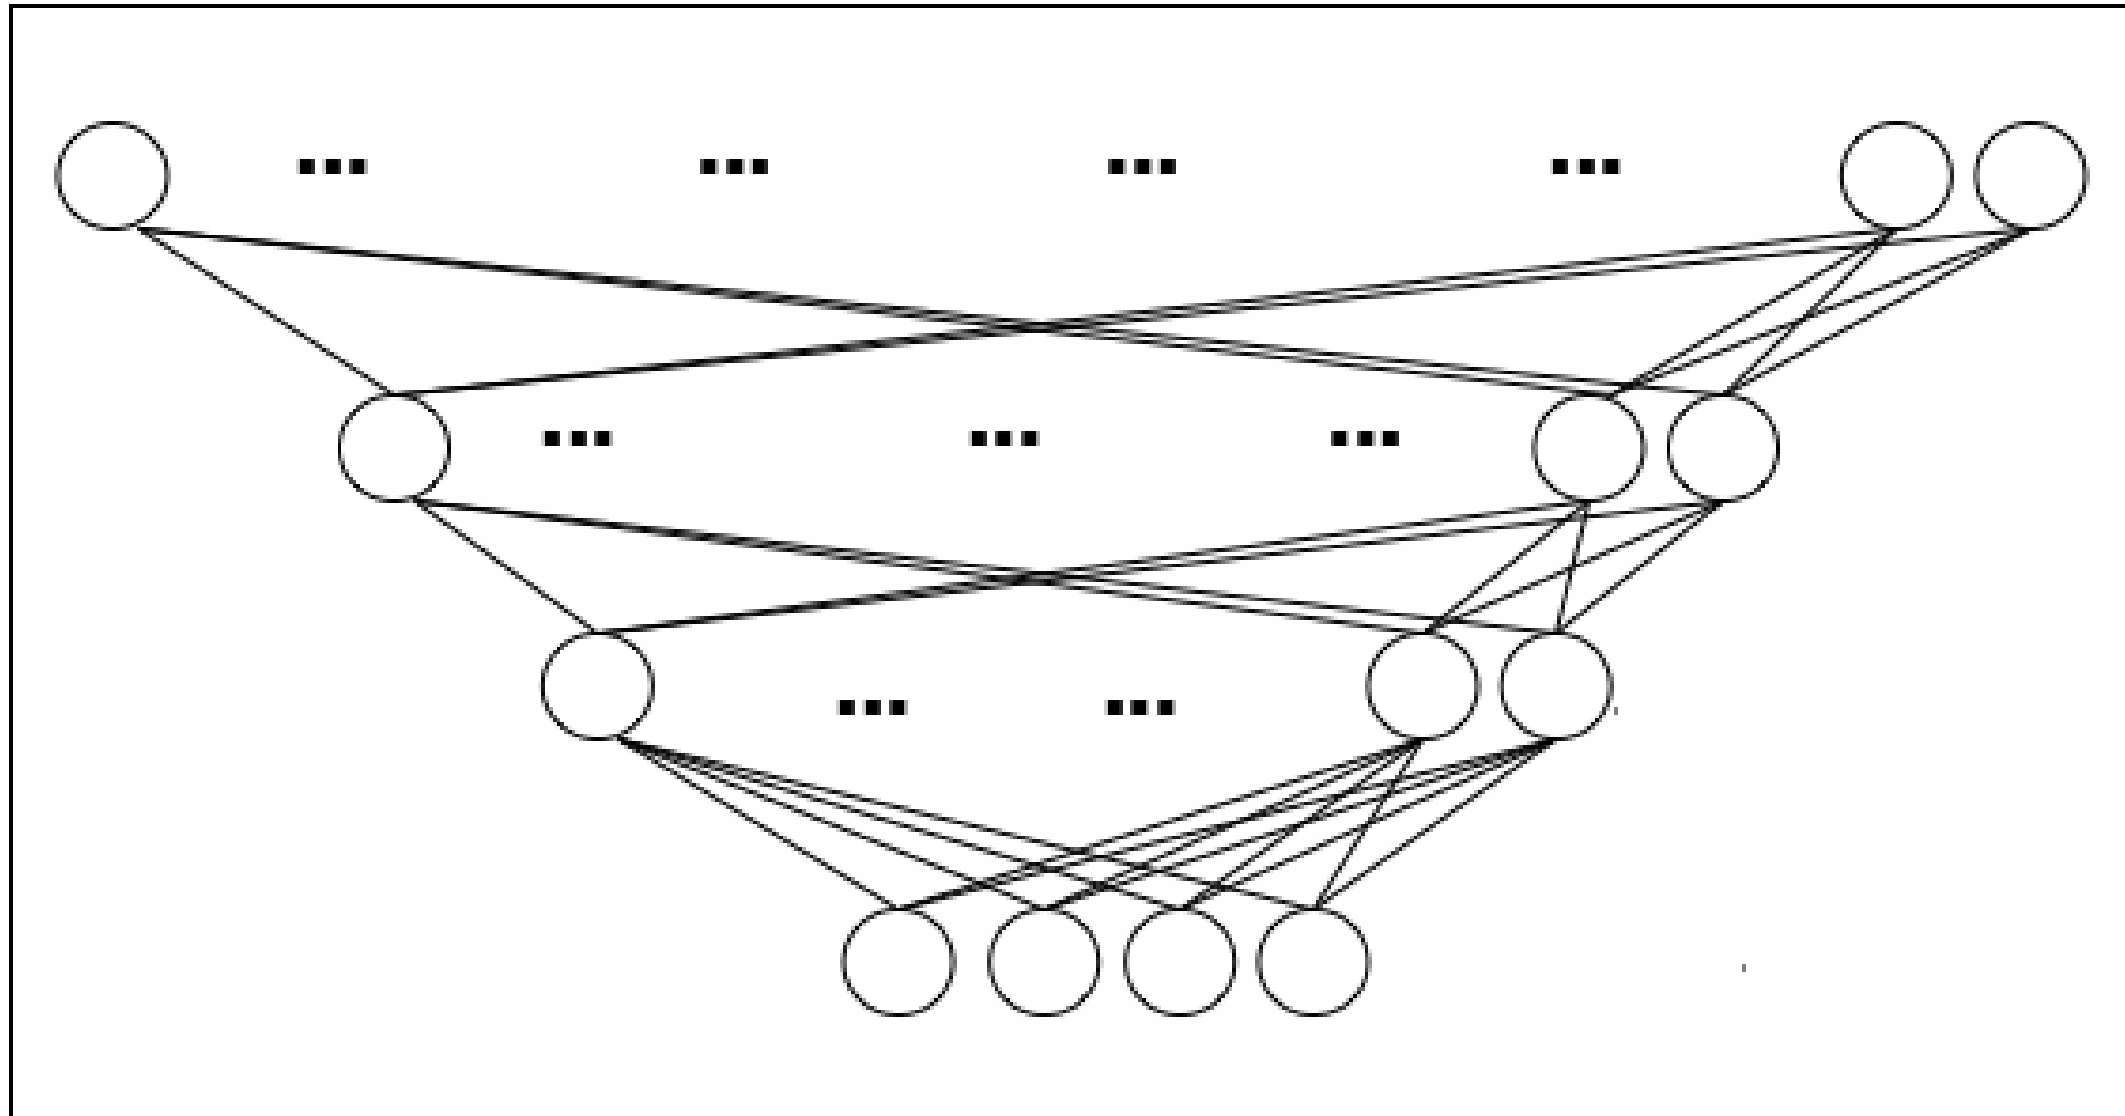
The sequential API

- Input layer

- Hidden layers

- Output layer

- Ordered in sequence

Building a sequential model

```python
# Import tensorflow
from tensorflow import keras

# Define a sequential model
model = keras.Sequential()
```

```python
# Define first hidden layer
model.add(keras.layers.Dense(16, activation='relu', input_shape=(28*28,)))
```

Building a sequential model

```python
# Define second hidden layer
model.add(keras.layers.Dense(8, activation='relu'))
```
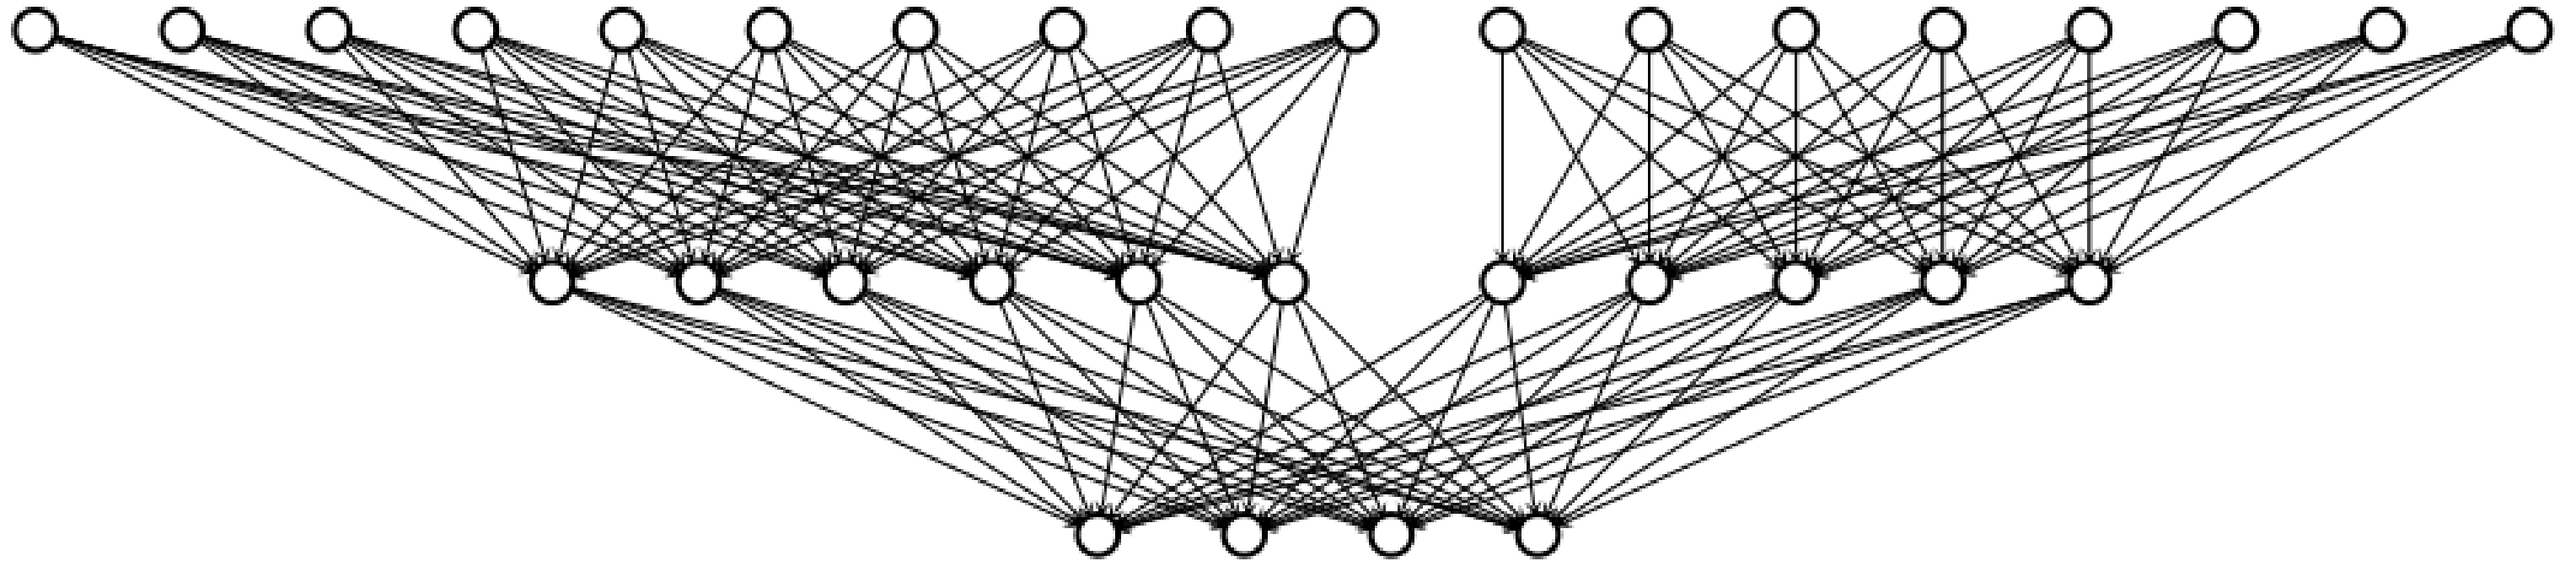
```python
# Define output layer
model.add(keras.layers.Dense(4, activation='softmax'))
```

```python
# Compile the model
model.compile('adam', loss='categorical_crossentropy')
```

```python
# Summarize the model
print(model.summary())
```

The functional API

INTRODUCTION TO TENSORFLOW IN PYTHON

Using the functional API

```python
# Import tensorflow
import tensorflow as tf


# Define model 1 input layer shape
model1_inputs = tf.keras.Input(shape=(28*28,))


# Define model 2 input layer shape
model2_inputs = tf.keras.Input(shape=(10,))
```

```python
# Define layer 1 for model 1
model1_layer1 = tf.keras.layers.Dense(12, activation='relu')(model1_inputs)


# Define layer 2 for model 1
model1_layer2 = tf.keras.layers.Dense(4, activation='softmax')(model1_layer1)
```

Using the functional API

```python
# Define layer 1 for model 2
model2_layer1 = tf.keras.layers.Dense(8, activation='relu')(model2_inputs)


# Define layer 2 for model 2
model2_layer2 = tf.keras.layers.Dense(4, activation='softmax')(model2_layer1)
```

```python
# Merge model 1 and model 2
merged = tf.keras.layers.add([model1_layer2, model2_layer2])
```

```python
# Define a functional model
model = tf.keras.Model(inputs=[model1_inputs, model2_inputs], outputs=merged)

# Compile the model
model.compile('adam', loss='categorical_crossentropy')
```

Let's practice!

datacamp

Training and validation with Keras

INTRODUCTION TO TENSORFLOW IN PYTHON

Isaiah Hull
Economist

Overview of training and evaluation

1. Load and clean data

2. Define model

3. Train and validate model

4. Evaluate model

How to train a model

```python
# Import tensorflow
import tensorflow as tf

# Define a sequential model
model = tf.keras.Sequential()
```

```python
# Define the hidden layer
model.add(tf.keras.layers.Dense(16, activation='relu', input_shape=(784,)))
```

```python
# Define the output layer
model.add(tf.keras.layers.Dense(4, activation='softmax'))
```

How to train a model

```python
# Compile model
model.compile('adam', loss='categorical_crossentropy')
```

```python
# Train model
model.fit(image_features, image_labels)
```
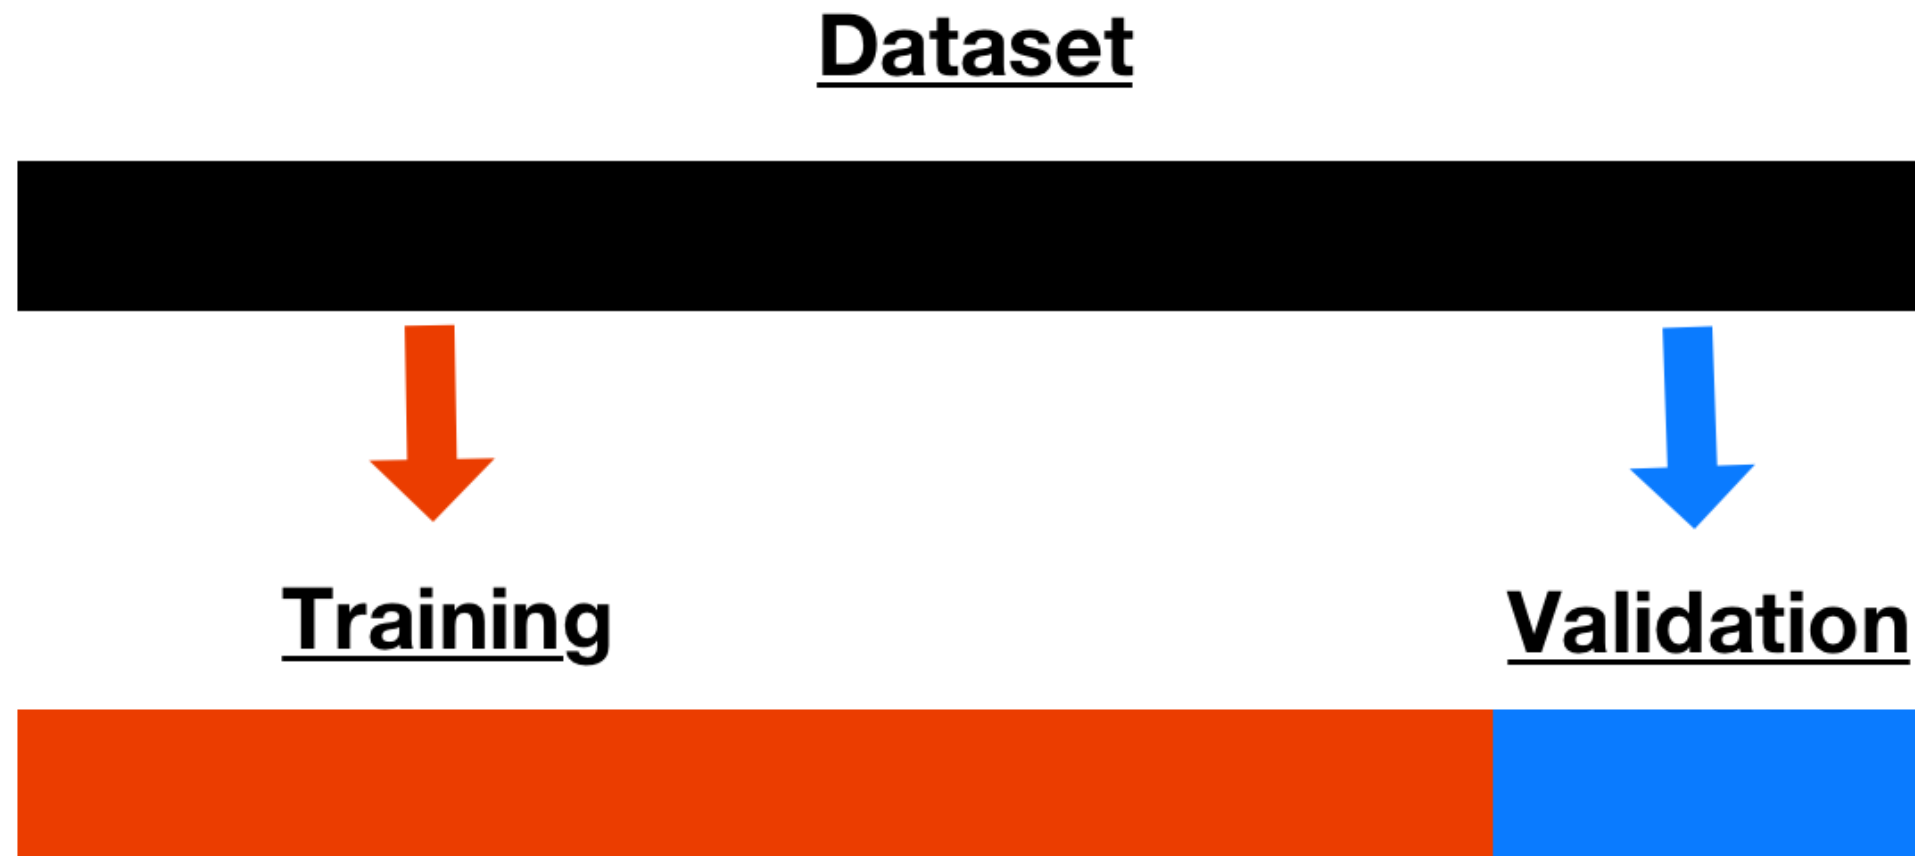
The fit() operation

- Required arguments
  - `features`

  - `labels`

- Many optional arguments
  - `batch_size`

  - `epochs`

  - `validation_split`

Batch size and epochs

Performing validation

Performing validation

```python
# Train model with validation split
model.fit(features, labels, epochs=10, validation_split=0.20)
```

Performing validation

```
Train on 1599 samples, validate on 400 samples
Epoch 1/10
1599/1599==============================] - 0s 159us/sample - loss: 1.2291 - val_loss: 1.0122
Epoch 2/10
1599/1599==============================] - 0s 60us/sample - loss: 0.8873 - val_loss: 0.7181
Epoch 3/10
1599/1599==============================] - 0s 61us/sample - loss: 0.6476 - val_loss: 0.5414
Epoch 4/10
1599/1599==============================] - 0s 58us/sample - loss: 0.4974 - val_loss: 0.4254
Epoch 5/10
1599/1599==============================] - 0s 57us/sample - loss: 0.3958 - val_loss: 0.3544
Epoch 6/10
1599/1599==============================] - 0s 62us/sample - loss: 0.3222 - val_loss: 0.2936
Epoch 7/10
1599/1599==============================] - 0s 58us/sample - loss: 0.2730 - val_loss: 0.2555
Epoch 8/10
1599/1599==============================] - 0s 56us/sample - loss: 0.2320 - val_loss: 0.2131
Epoch 9/10
1599/1599==============================] - 0s 59us/sample - loss: 0.1957 - val_loss: 0.1843
Epoch 10/10
1599/1599==============================] - 0s 55us/sample - loss: 0.1663 - val_loss: 0.1657
```

Changing the metric

```python
# Recomile the model with the accuracy metric
model.compile('adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```python
# Train model with validation split
model.fit(features, labels, epochs=10, validation_split=0.20)
```

Changing the metric

```
Train on 1599 samples, validate on 400 samples
Epoch 1/10
1599/1599==========================] - 0s 174us/sample - loss: 1.2956 - acc: 0.4196 - val_loss: 1.1189 - val_acc: 0.5075
Epoch 2/10
1599/1599==========================] - 0s 59us/sample - loss: 0.9356 - acc: 0.7949 - val_loss: 0.7843 - val_acc: 0.8225
Epoch 3/10
1599/1599==========================] - 0s 59us/sample - loss: 0.6657 - acc: 0.9037 - val_loss: 0.5588 - val_acc: 0.8925
Epoch 4/10
1599/1599==========================] - 0s 58us/sample - loss: 0.4898 - acc: 0.9206 - val_loss: 0.4220 - val_acc: 0.9175
Epoch 5/10
1599/1599==========================] - 0s 59us/sample - loss: 0.3734 - acc: 0.9681 - val_loss: 0.3319 - val_acc: 0.9825
Epoch 6/10
1599/1599==========================] - 0s 61us/sample - loss: 0.2975 - acc: 0.9762 - val_loss: 0.2907 - val_acc: 0.9075
Epoch 7/10
1599/1599==========================] - 0s 60us/sample - loss: 0.2414 - acc: 0.9731 - val_loss: 0.2276 - val_acc: 0.9550
Epoch 8/10
1599/1599==========================] - 0s 61us/sample - loss: 0.1912 - acc: 0.9887 - val_loss: 0.2026 - val_acc: 0.9525
Epoch 9/10
1599/1599==========================] - 0s 59us/sample - loss: 0.1649 - acc: 0.9862 - val_loss: 0.1684 - val_acc: 0.9675
Epoch 10/10
1599/1599==========================] - 0s 58us/sample - loss: 0.1390 - acc: 0.9912 - val_loss: 0.1374 - val_acc: 0.9825
```
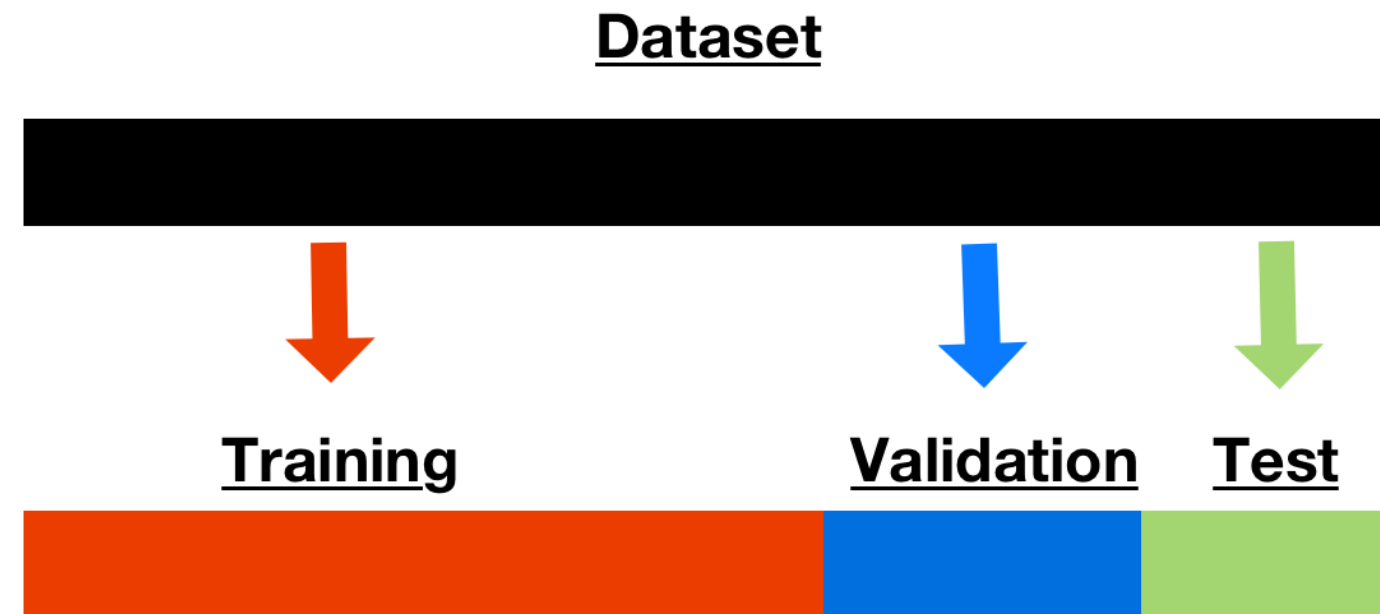
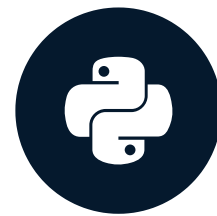# The evaluation() operation

Dataset


```
# Evaluate the test set
model.evaluate(test)
```

Let's practice!

INTRODUCTION TO TENSORFLOW IN PYTHON

Training models with the Estimators API

INTRODUCTION TO TENSORFLOW IN PYTHON

Isaiah Hull
Economist

What is the Estimators API?

- High level submodule

- Less flexible

- Enforces best practices

- Faster deployment

- Many premade models

High-Level
TensorFlow APIs

| Estimators |

Mid-Level
TensorFlow APIs

| Layers | Datasets | Metrics |

Low-level
TensorFlow APIs

| Python |

[1] Image taken from https://www.tensorflow.org/guide/premade_estimators

Model specification and training

1. Define feature columns

2. Load and transform data

3. Define an estimator

4. Apply train operation

Defining feature columns

```python
# Import tensorflow under its standard alias
import tensorflow as tf


# Define a numeric feature column
size = tf.feature_column.numeric_column("size")
```

```python
# Define a categorical feature column
rooms = tf.feature_column.categorical_column_with_vocabulary_list("rooms",\
["1", "2", "3", "4", "5"])
```

Defining feature columns

```python
# Create feature column list
features_list = [size, rooms]
```

```python
# Define a matrix feature column
features_list = [tf.feature_column.numeric_column('image', shape=(784,))]
```

Loading and transforming data

```python
# Define input data function
def input_fn():
    # Define feature dictionary
    features = {"size": [1340, 1690, 2720], "rooms": [1, 3, 4]}
    # Define labels
    labels = [221900, 538000, 180000]
    return features, labels
```

Define and train a regression estimator

```python
# Define a deep neural network regression
model0 = tf.estimator.DNNRegressor(feature_columns=feature_list,\
    hidden_units=[10, 6, 6, 3])


# Train the regression model
model0.train(input_fn, steps=20)
```

Define and train a deep neural network

```python
# Define a deep neural network classifier
model1 = tf.estimator.DNNClassifier(feature_columns=feature_list,\
    hidden_units=[32, 16, 8], n_classes=4)


# Train the classifier
model1.train(input_fn, steps=20)
```

- **https://www.tensorflow.org/guide/estimators**

Let's practice!

INTRODUCTION TO TENSORFLOW IN PYTHON

Congratulations!

INTRODUCTION TO TENSORFLOW IN PYTHON

Isaiah Hull
Economist

What you learned

- **Chapter 1**
  - Low-level, basic, and advanced operations

  - Graph-based computation

  - Gradient computation and optimization

- **Chapter 2**
  - Data loading and transformation

  - Predefined and custom loss functions

  - Linear models and batch training

What you learned

- **Chapter 3**
  - Dense neural network layers
  - Activation functions
  - Optimization algorithms
  - Training neural networks

- **Chapter 4**
  - Neural networks in Keras
  - Training and validation
  - The Estimators API

TensorFlow extensions

- **TensorFlow Hub**
  - Pretrained models
  - Transfer learning

- **TensorFlow Probability**
  - More statistical distributions
  - Trainable distributions
  - Extended set of optimizers

[1] Screenshot from https://tfhub.dev.
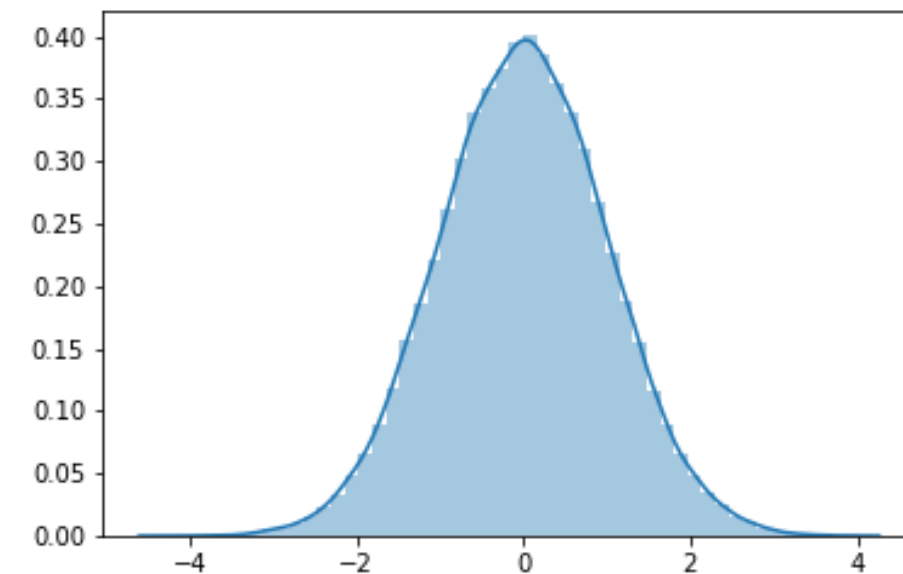
TensorFlow 2.0

- TensorFlow 2.0
  - `eager_execution()`

  - Tighter `keras` integration

  - `Estimators`

  - `function()`
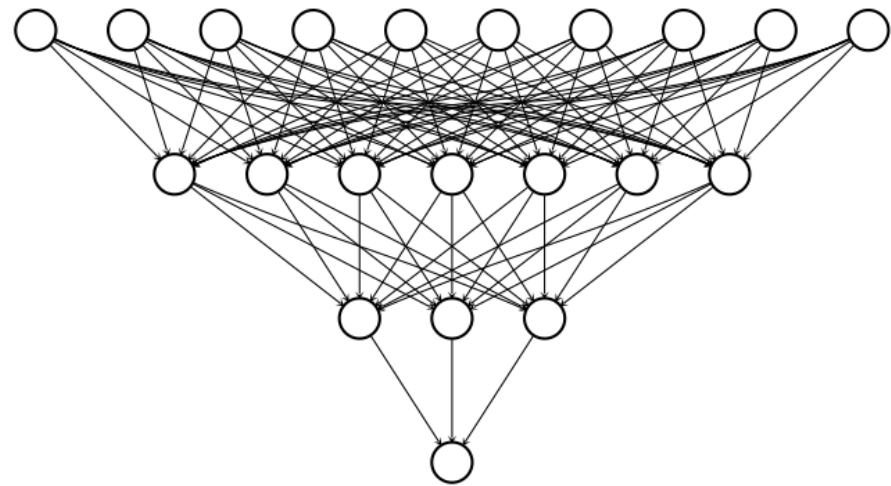
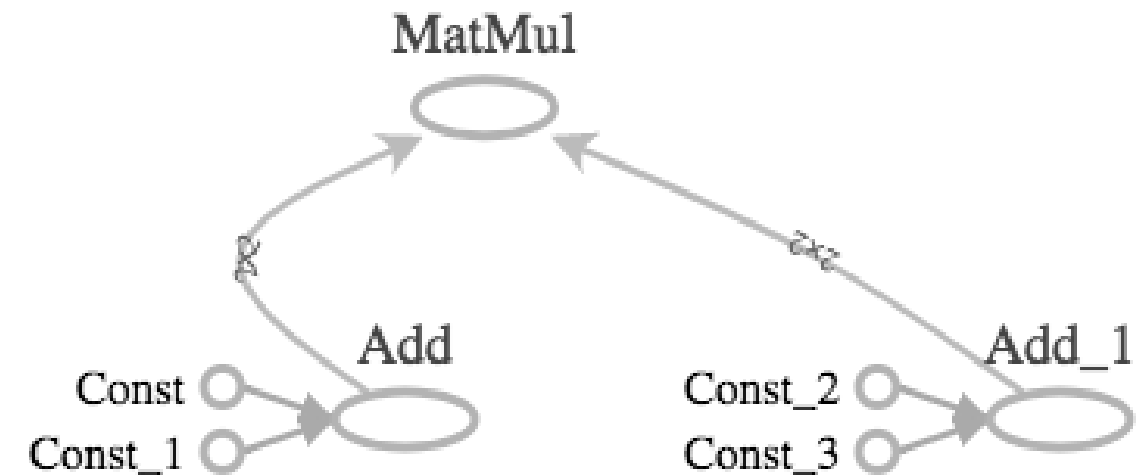|	
High-Level TensorFlow APIs	Estimators
Mid-Level TensorFlow APIs	Layers · Datasets · Metrics
Low-level TensorFlow APIs	Python

[1] Screenshot taken from https://www.tensorflow.org/guide/premade_estimators

Congratulations!

INTRODUCTION TO TENSORFLOW IN PYTHON

datacamp